

# Fault-Tolerant and Reproducible Input Processing for Machine Learning

Master Thesis

Dan Kluser supervised by Dan Graur and Prof. Ana Klimovic

# ML Input Data Processing

- More data is a key factor in Deep Learning's success
- Data is expensive
- Augmenting dataset needs compute intensive online processing
- Often pre-processing becomes a bottleneck
  - 20% of jobs spend 1/3 in input pipeline
  - 30% of total compute time for pre-processing

# Alleviating Bottleneck by Scaling Out

- Industry agrees: Both Meta and Google scale out preprocessing using disaggregated service
- Right-size resources for training and pre-processing independently
- Exploit synergies between jobs by having a fleet-wide view within the service

# Issues when Scaling Out

## 1) Fault-Tolerance

As more nodes are involved in a job failure probabilities become non-negligible

## 2) Reproducibility

Distributing the job introduces additional sources of non-determinism

*(DRR is joint work with Zak; presented by him)*

# Issues when Scaling Out

**This presentation**

## 1) Fault-Tolerance

As more nodes are involved in a job failure probabilities become non-negligible

# Motivation for Fault-Tolerance

- Only 35.9% ML jobs at Meta did not experience any fault [over course of a week]
- Use of transient cloud resources (i.e. SpotVMs)
  - At least 60% cheaper, but may be preempted at any time
- Live migration of nodes in datacenters

# Requirements for Fault-Tolerant Processing

1. Disaggregated Distributed Pre-Processing
2. Correctness (Exactly-Once)
3. Performance (Bounded overhead; recover some progress)
4. Compatibility
5. Reproducibility

# Platform to build on?

- Know of two disaggregated distributed preprocessing services: tf.data service and Meta's DPP
- Use Cachew (built on tf.data service) as it is open source
- Relieves ML users from the burden of managing compute, memory and storage infrastructure for ML input preprocessing
- Builds on top of tf.data service and provides:
  - Distributed disaggregated multi-tenant input pre-processing (Req1)
  - Supports tf.data pipelines (Req4)
  - Autocaching and autoscaling




# Mitigating Worker Failure Impact on Accuracy

How often should a client see an example per epoch?

- No-guarantees, except for the same epoch size
  - See figure on next-slide, accuracy suffers, double-digit drop
- At-most-once
  - Could lead to some examples never being seen
  - Problematic with highly class-imbalanced datasets
- Exactly-once
  - No possibility of model performance degrading due to “invisible” pre-processing failures

# CheckFreq Accuracy Drop

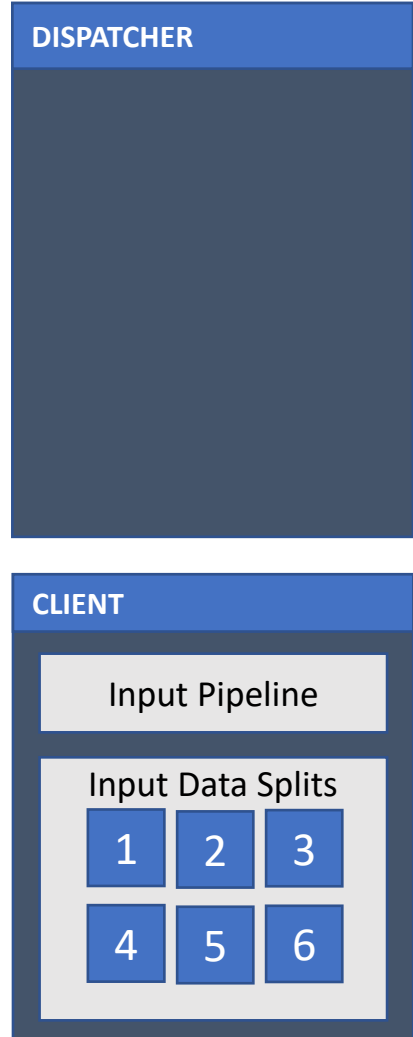


## Takeaway

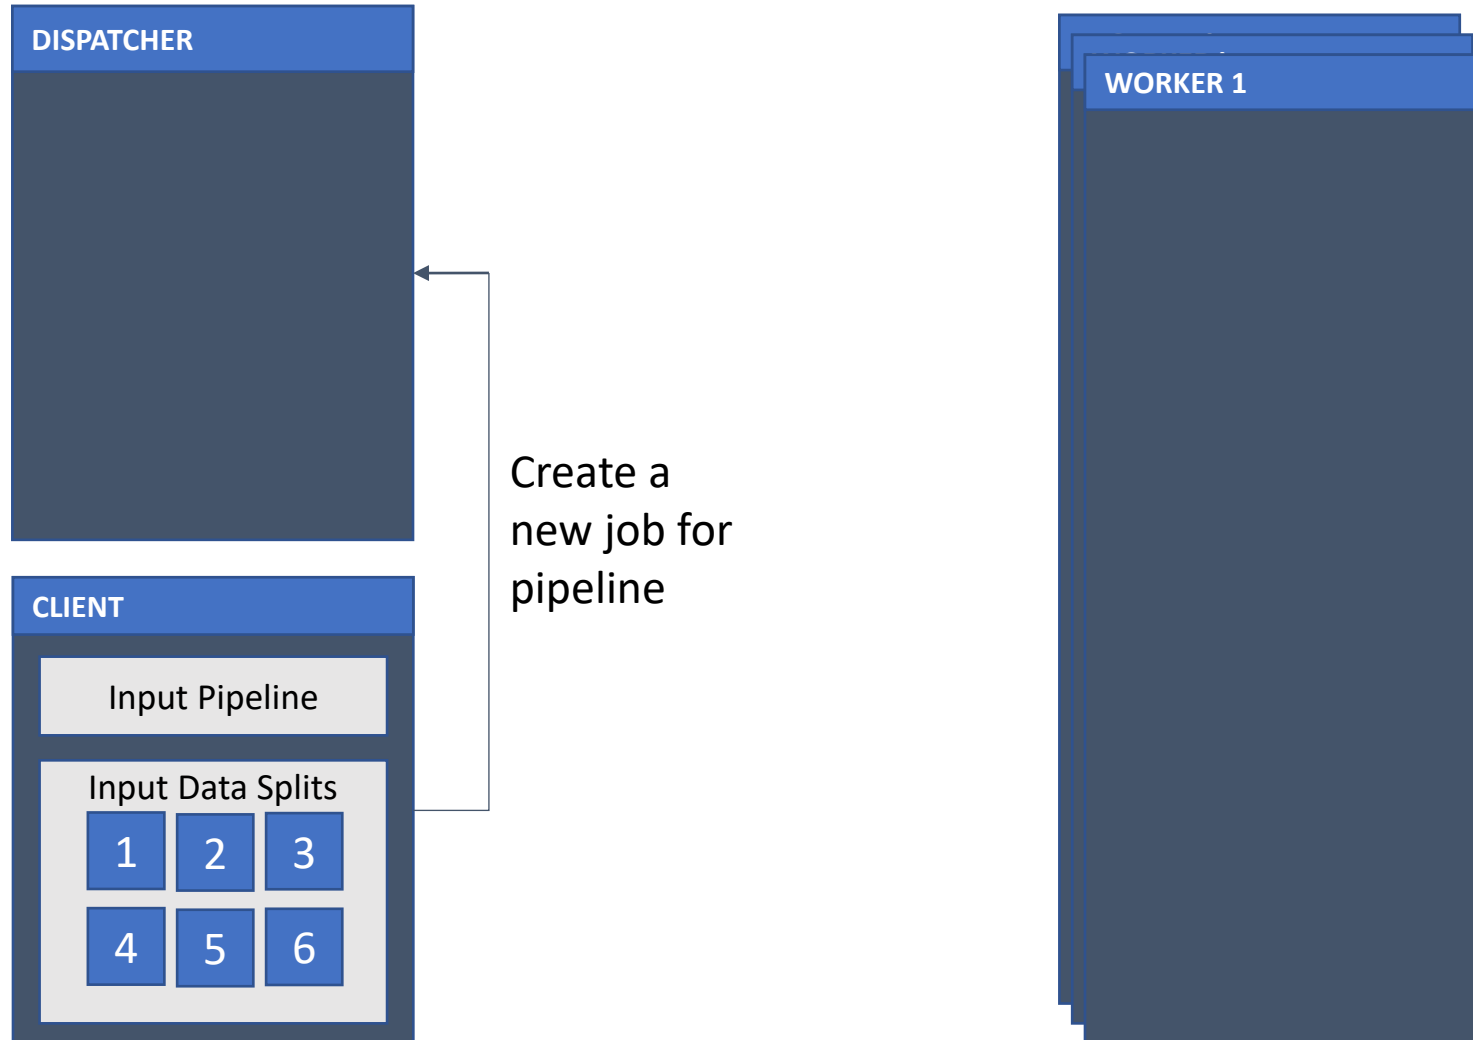
Duplicated / missing examples lead to a 13% drop in accuracy

Mohan, J., Phanishayee, A., & Chidambaram, V. (2021, February). CheckFreq: Frequent, Fine-Grained DNN Checkpointing. In Proceedings of the USENIX Conference on File and Storage Technologies (FAST 2021).

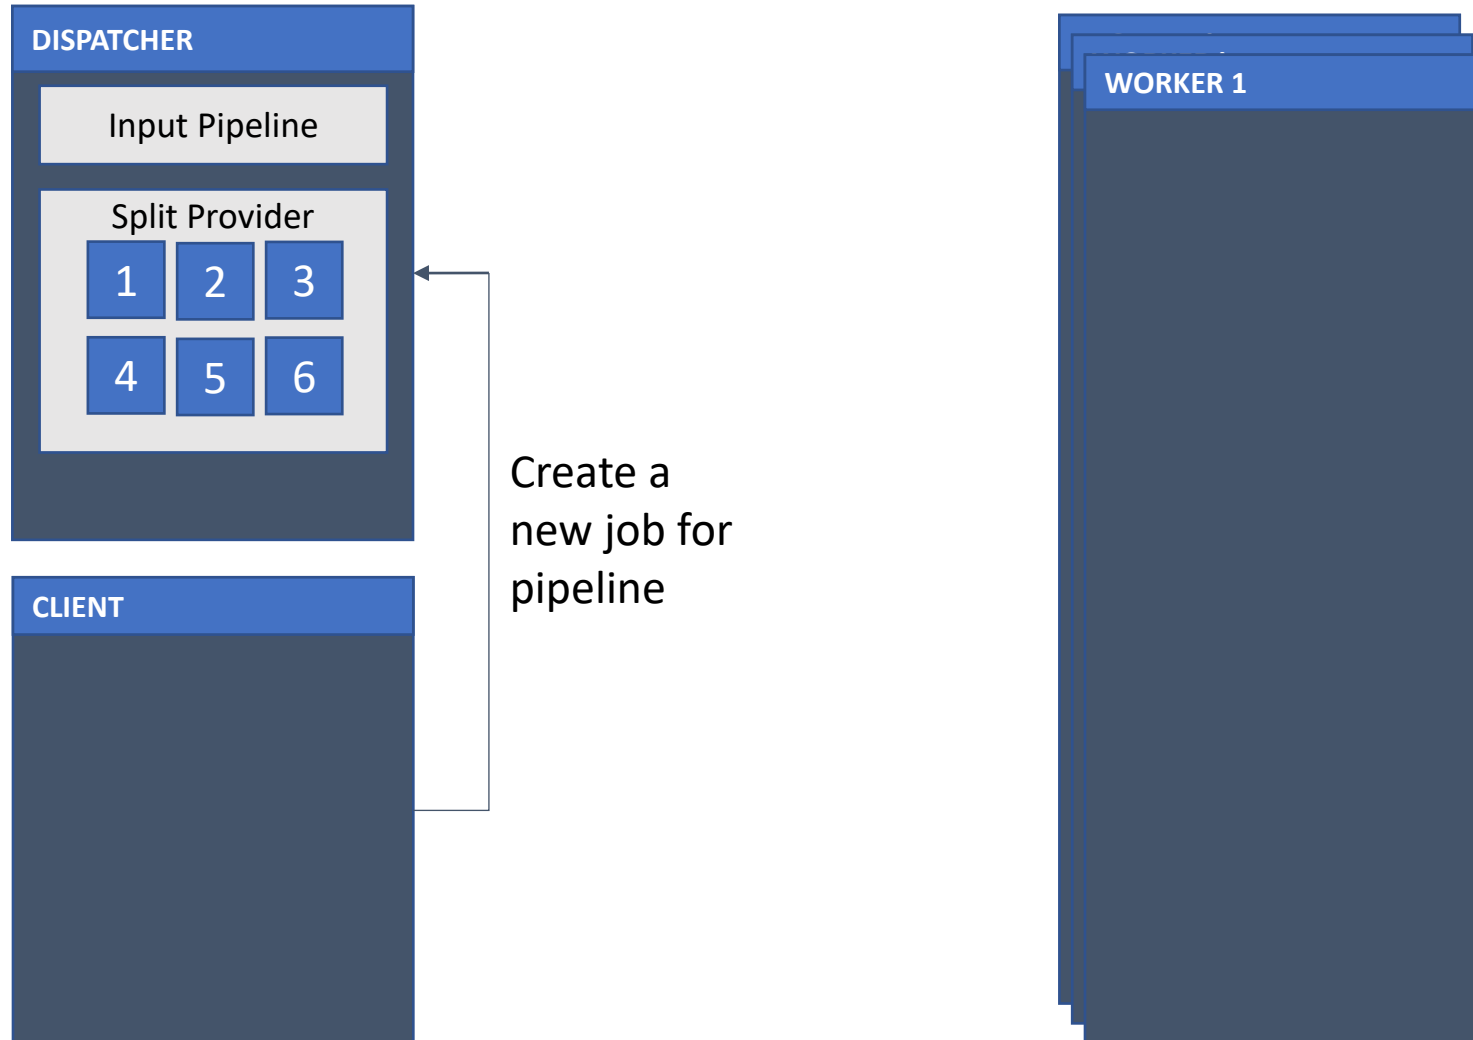
# Cachew Illustrated



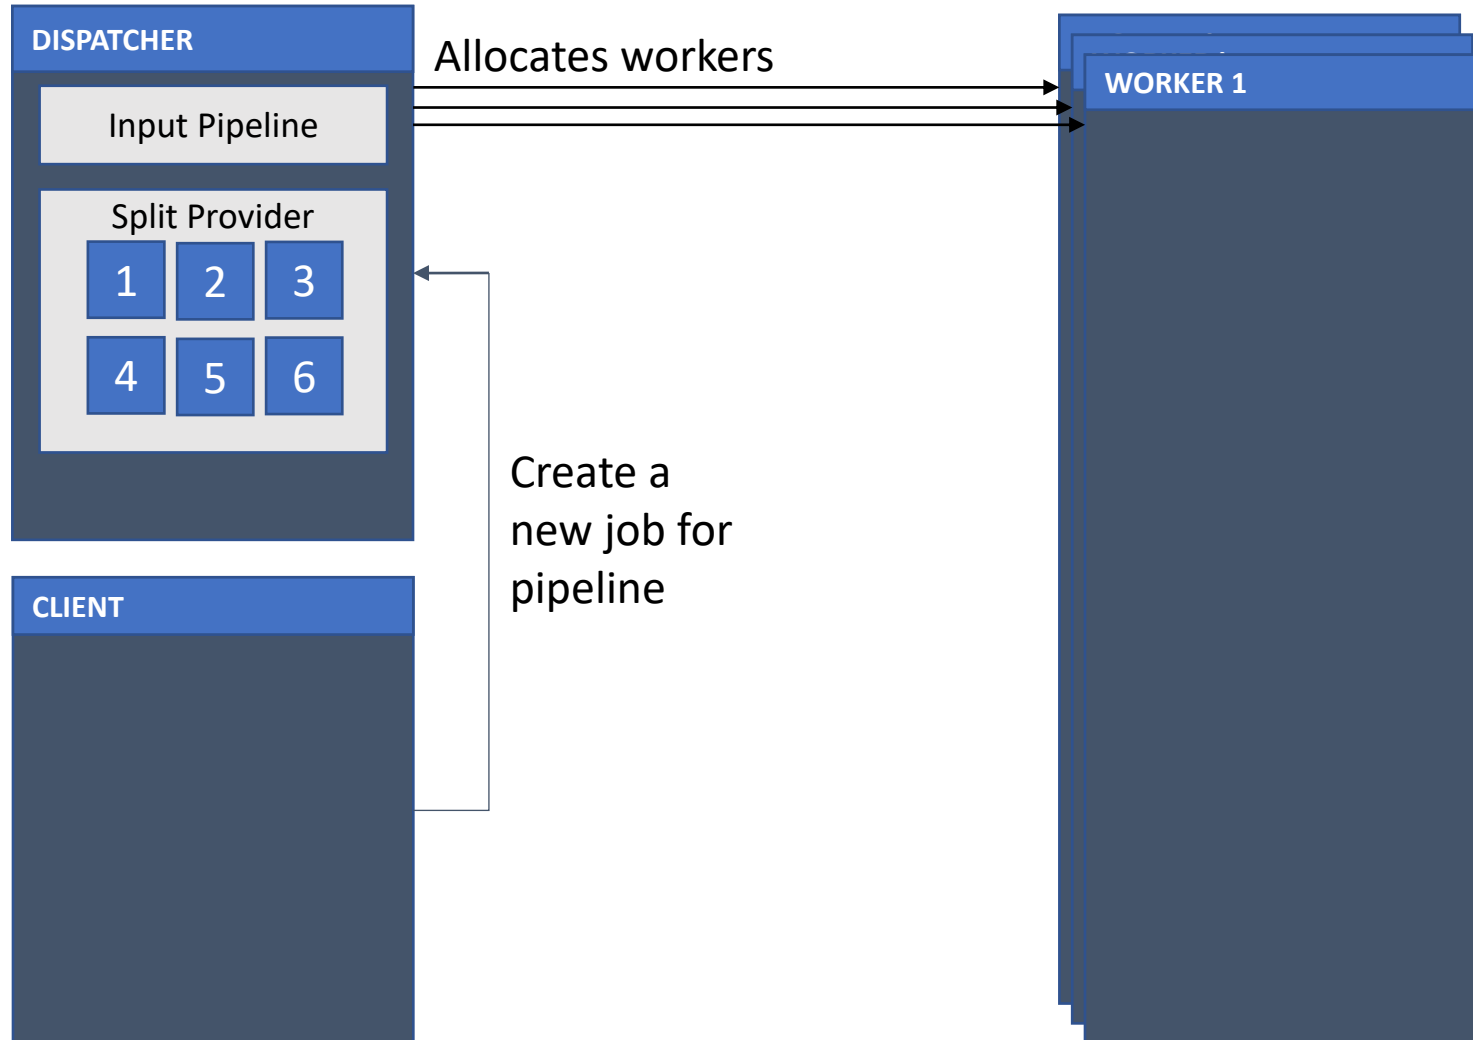
# Cachew Illustrated



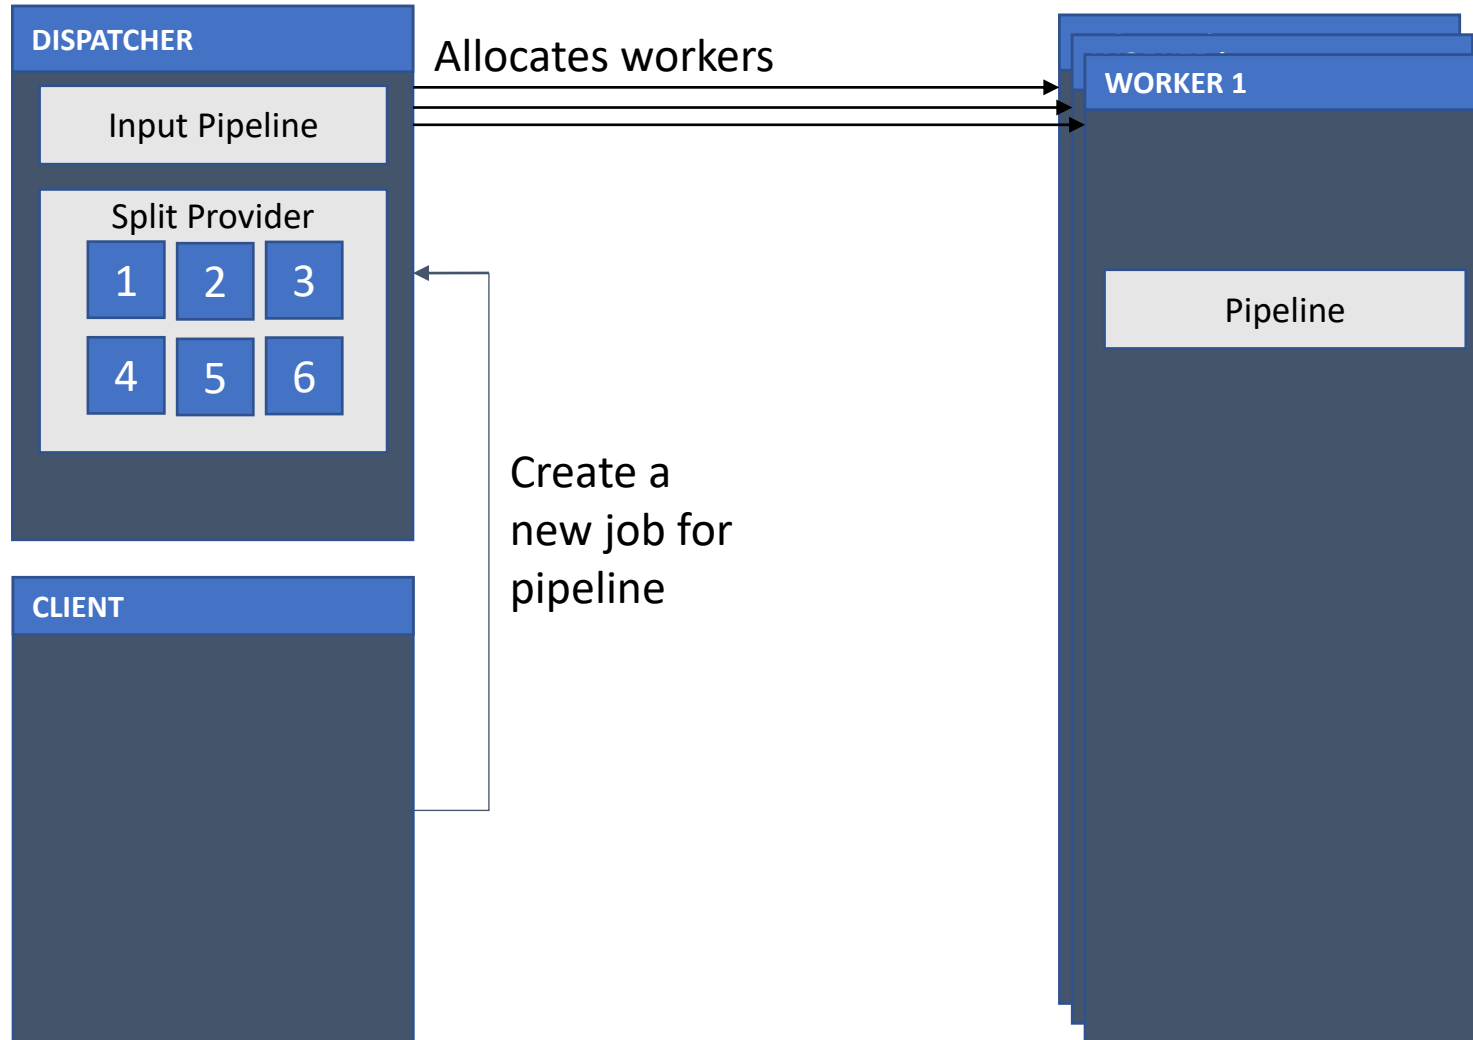
# Cachew Illustrated



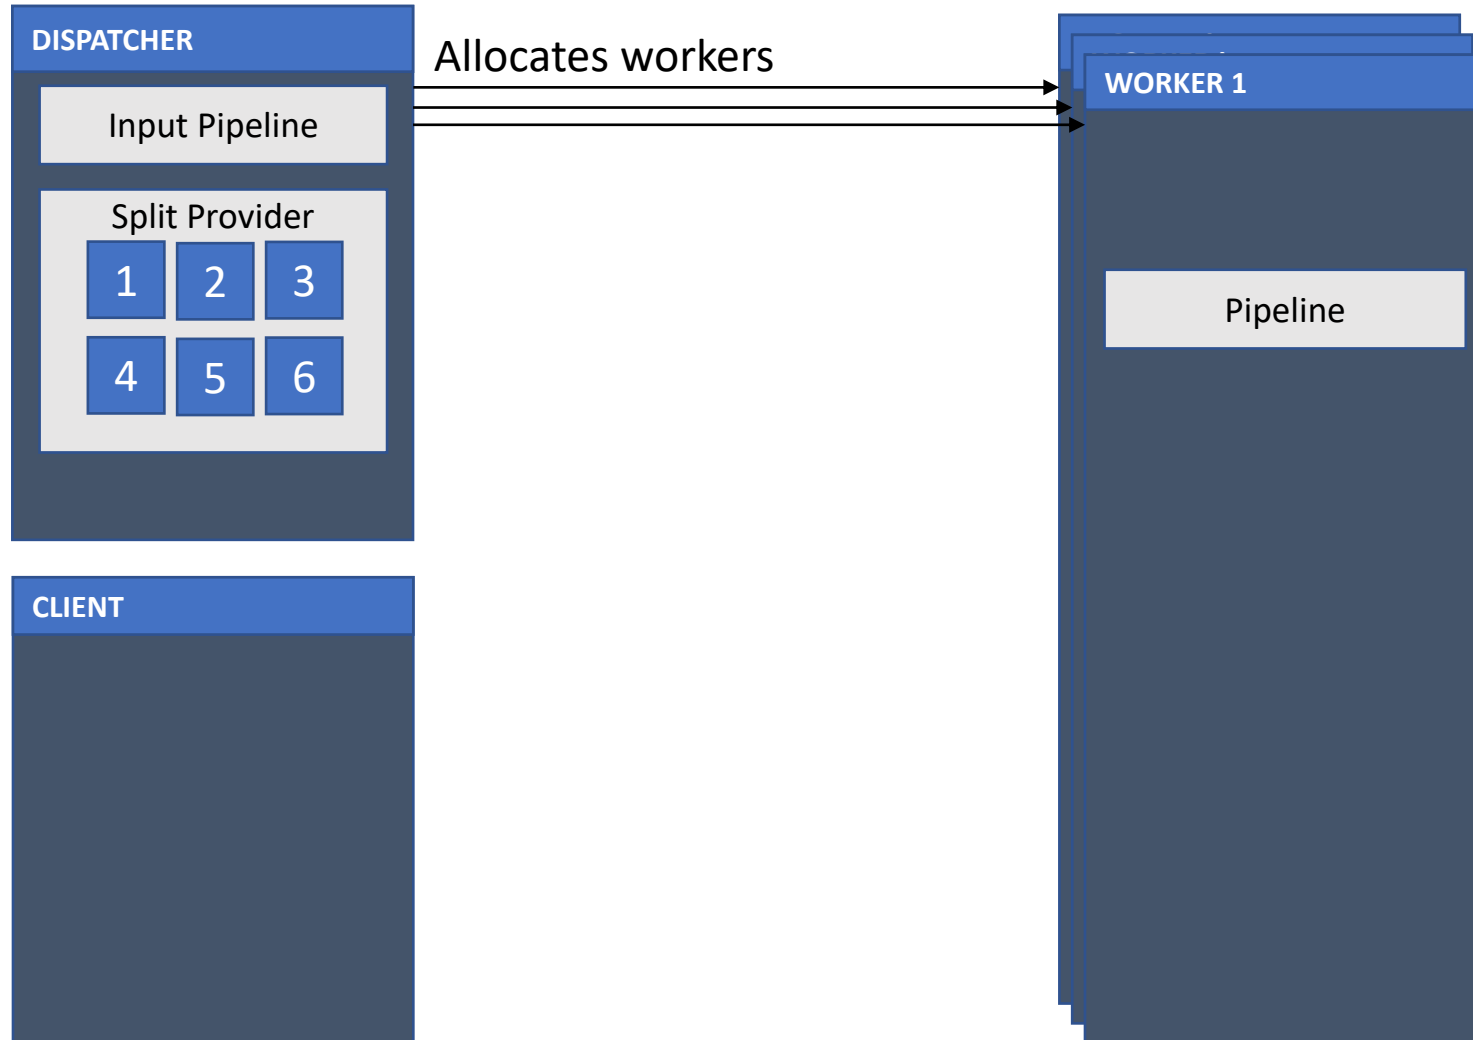
# Cachew Illustrated



# Cachew Illustrated

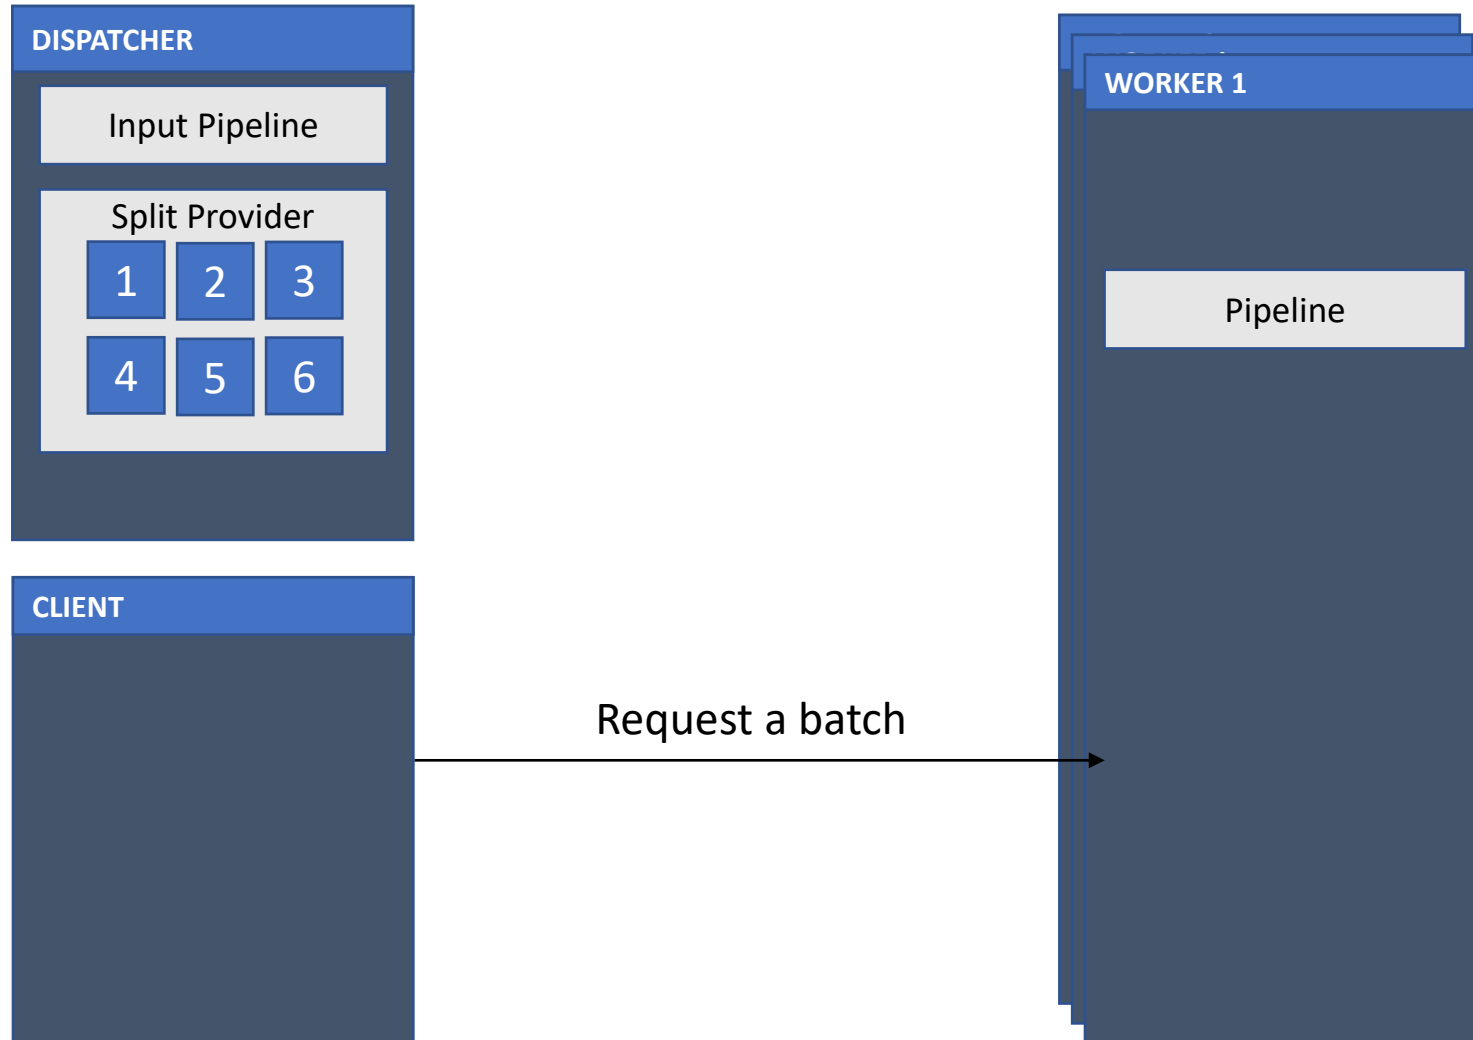


# Cachew Illustrated

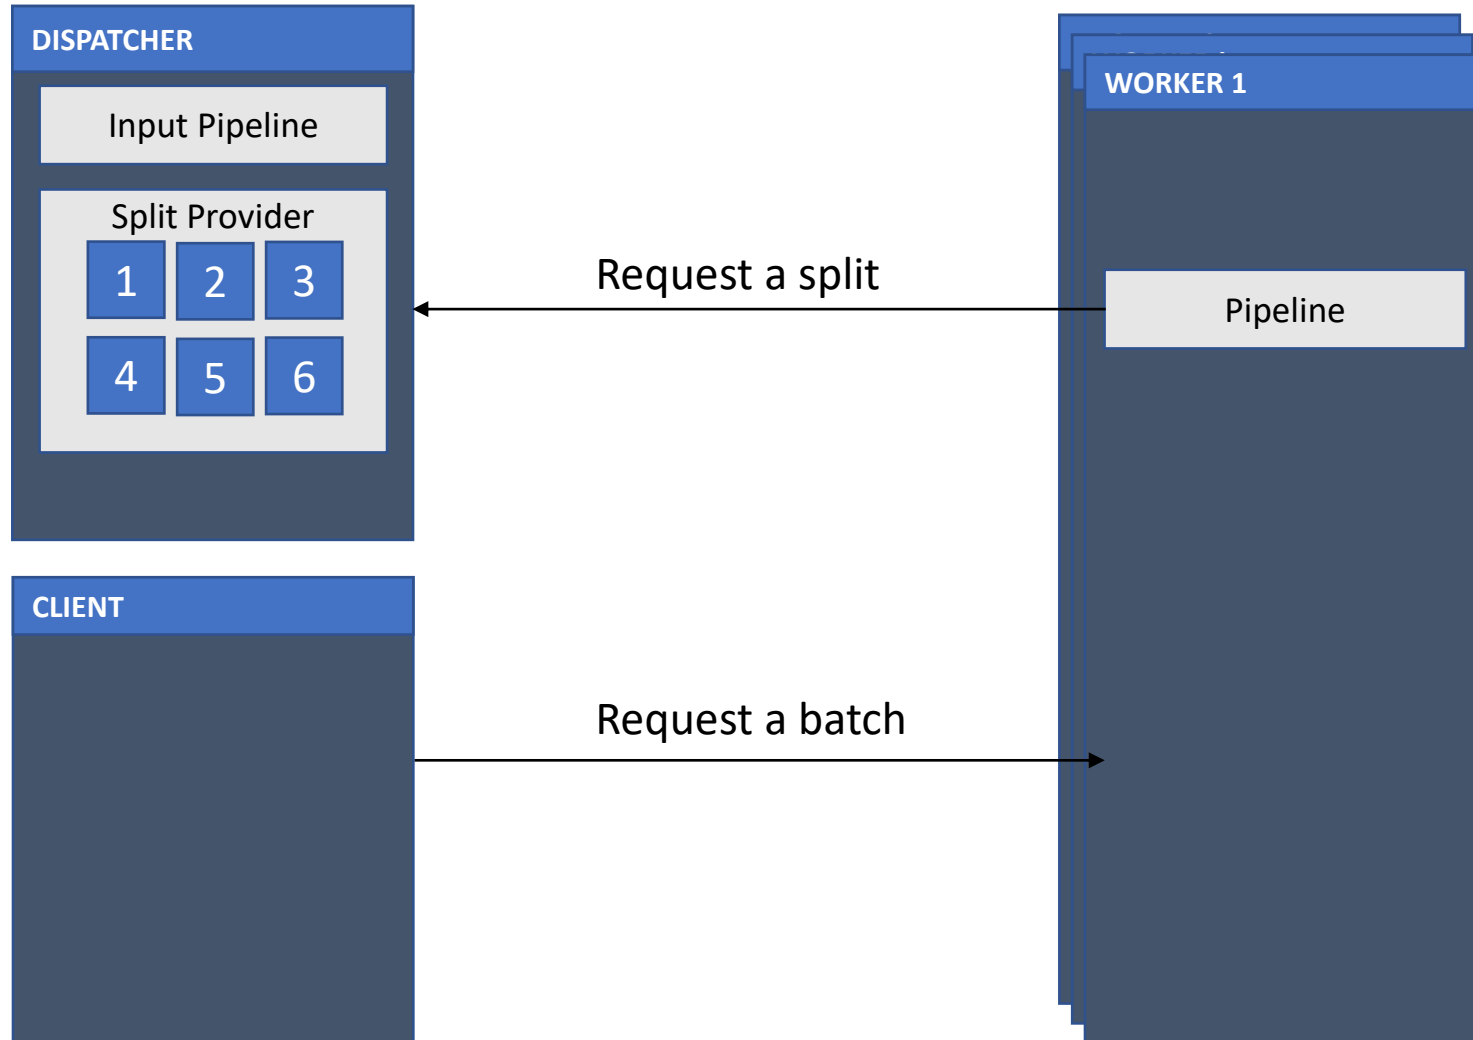




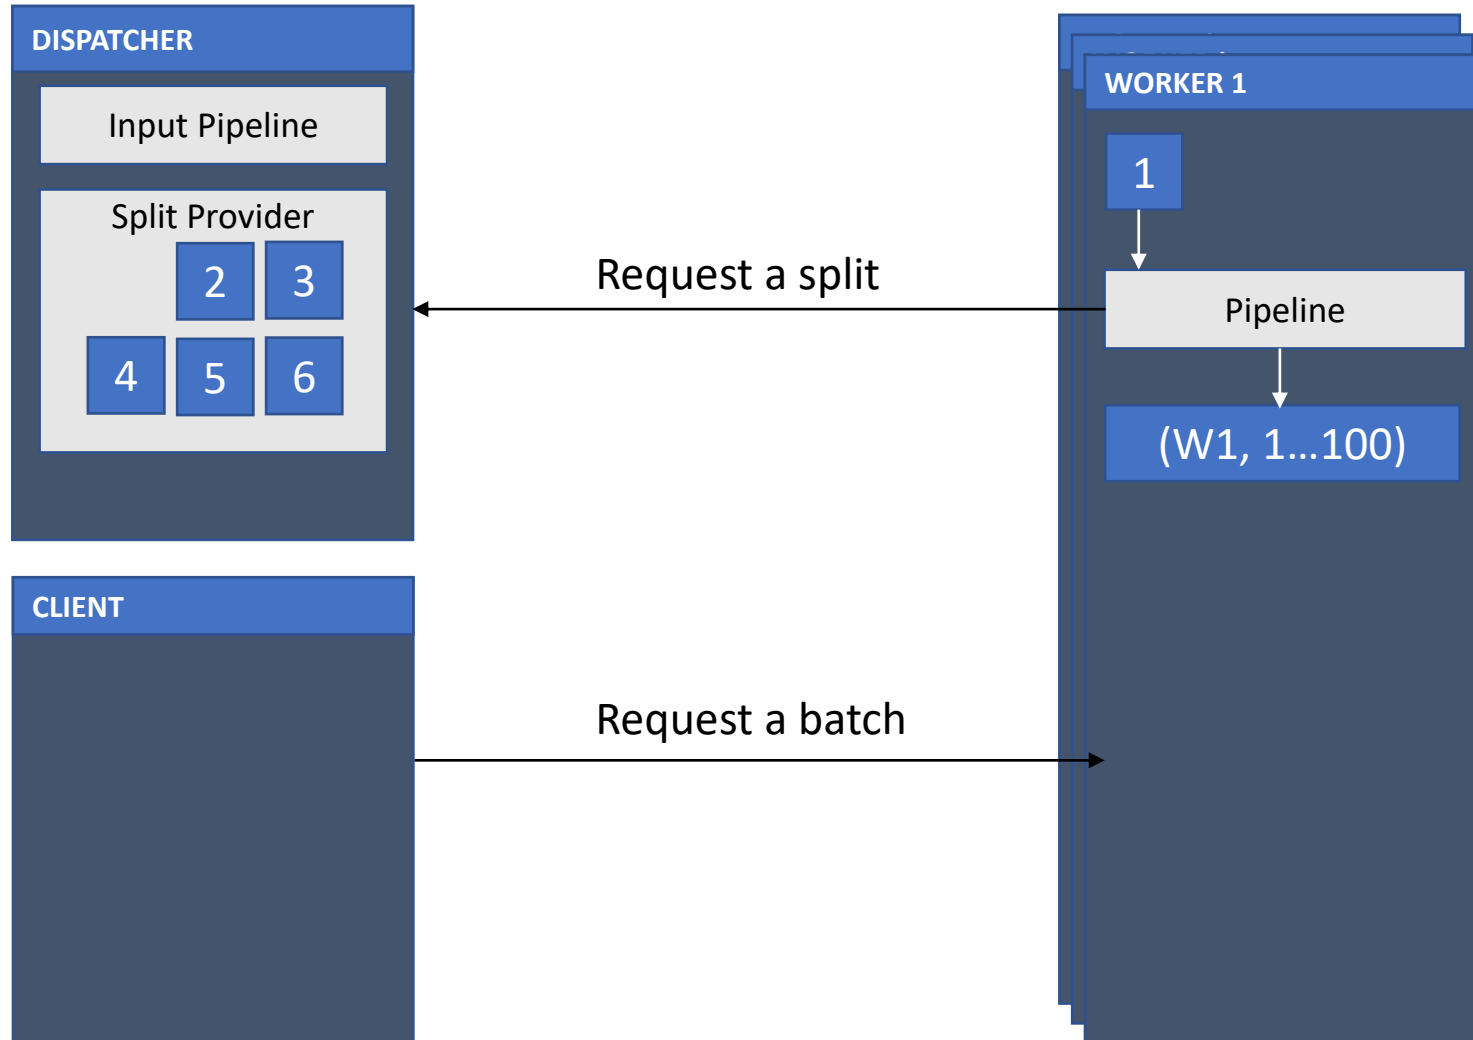
# Cachew Illustrated



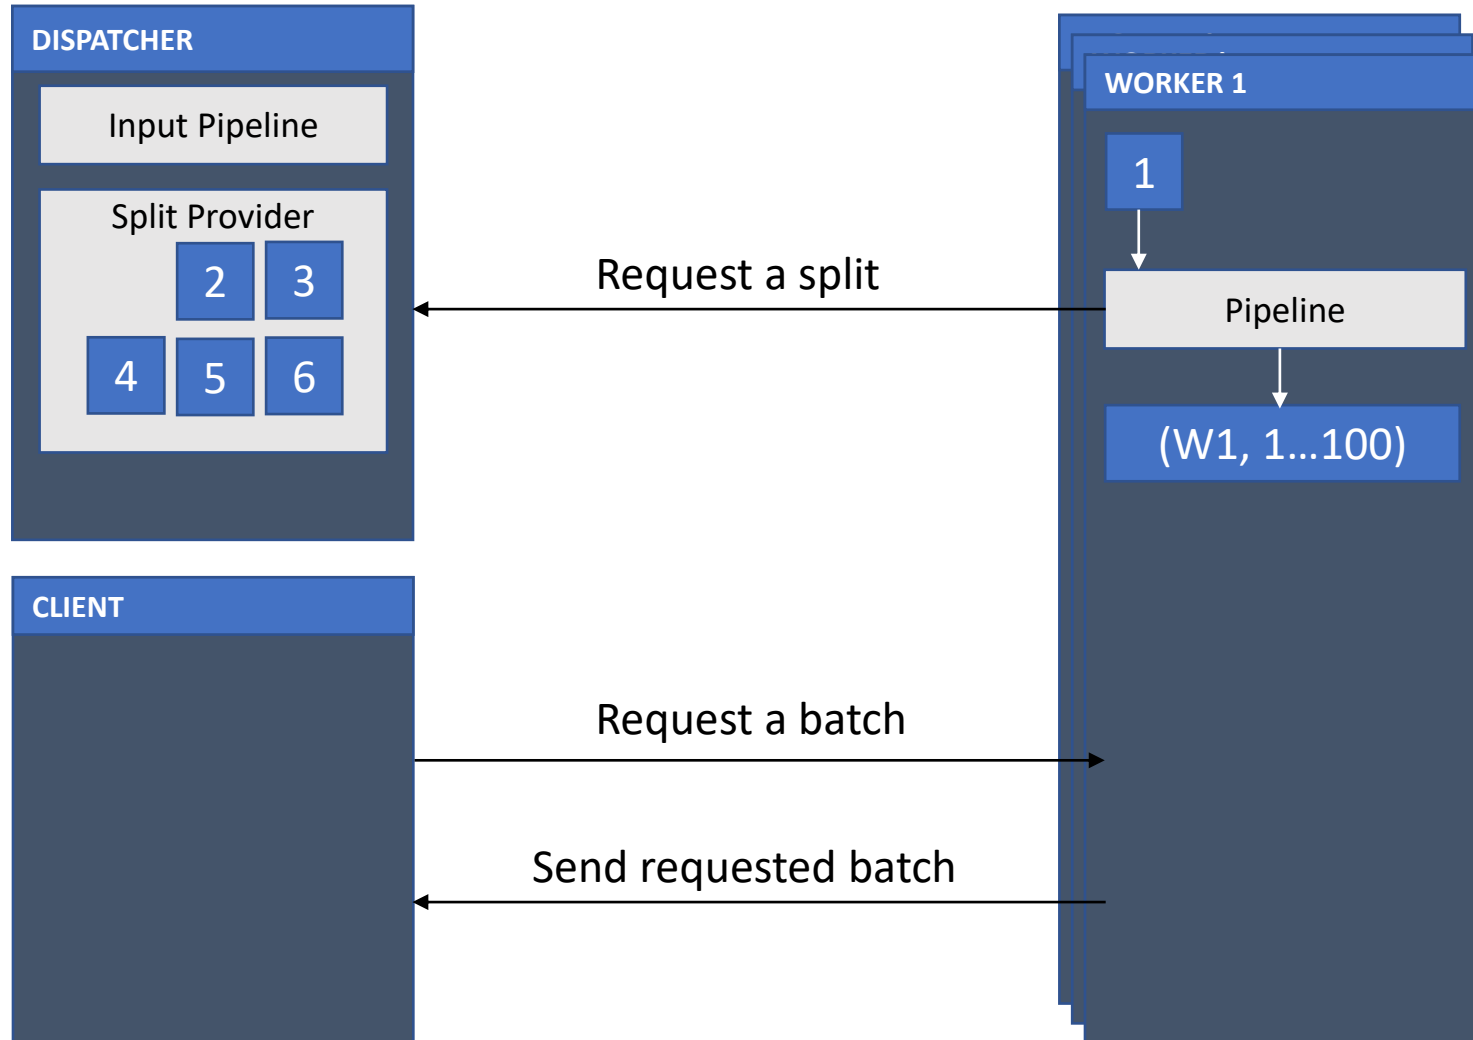
# Cachew Illustrated



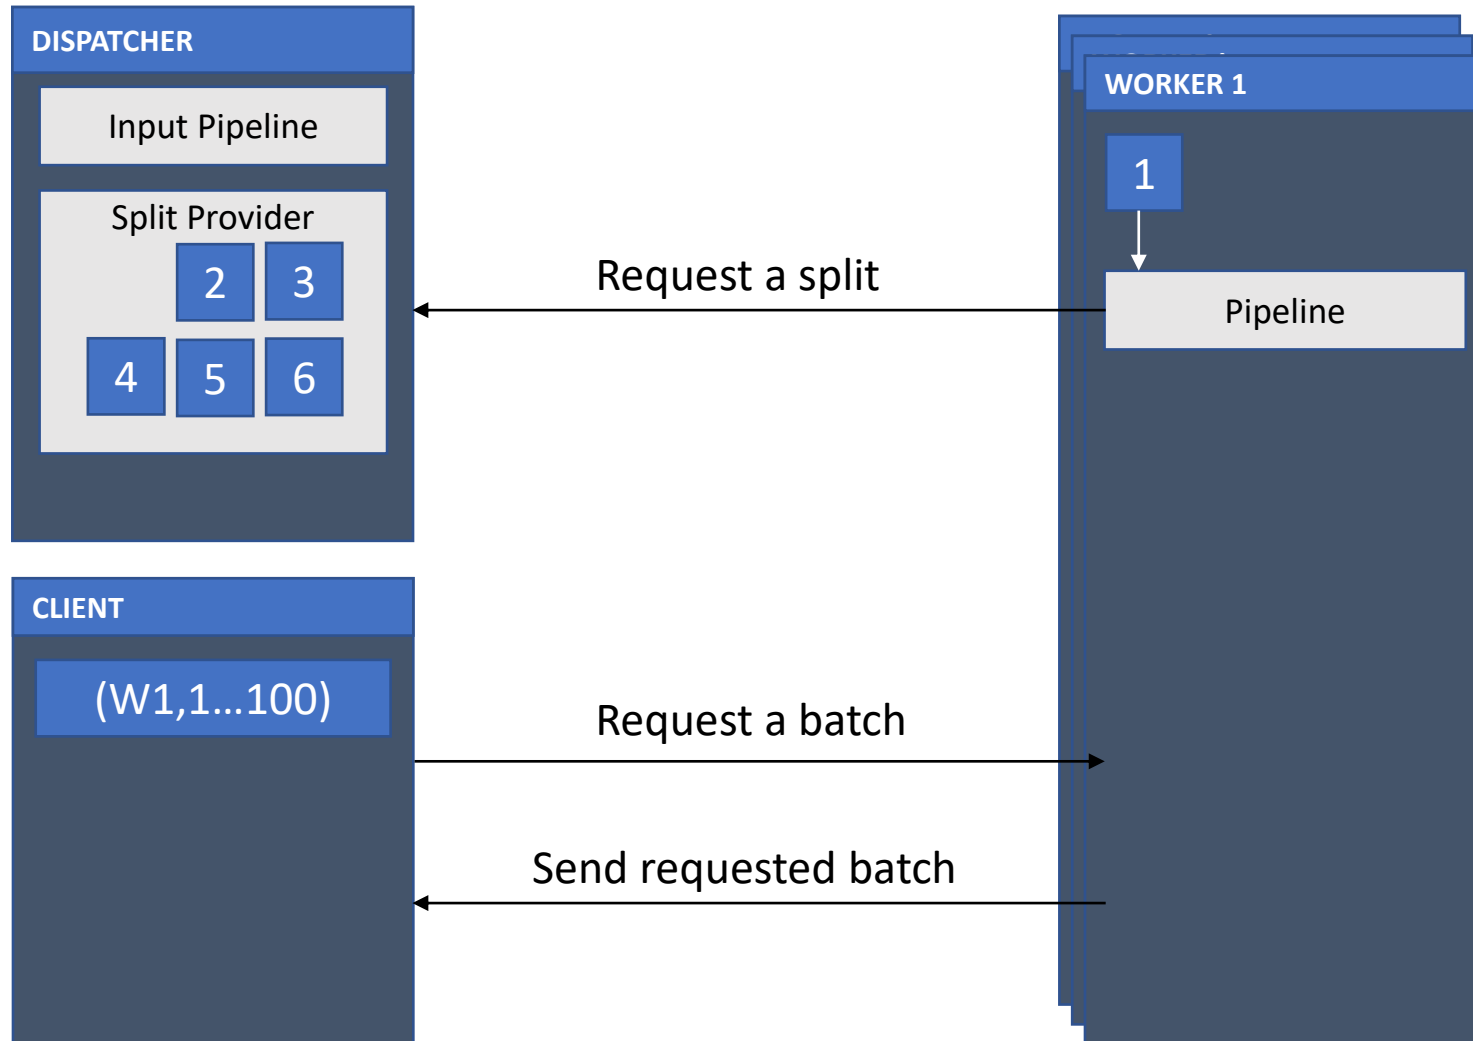
# Cachew Illustrated




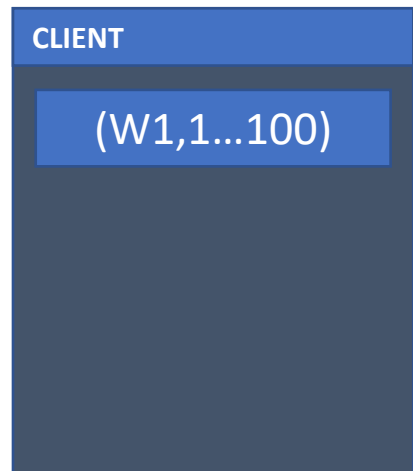
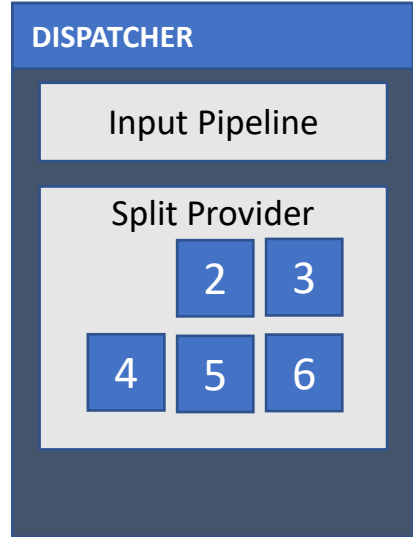
# Cachew Illustrated



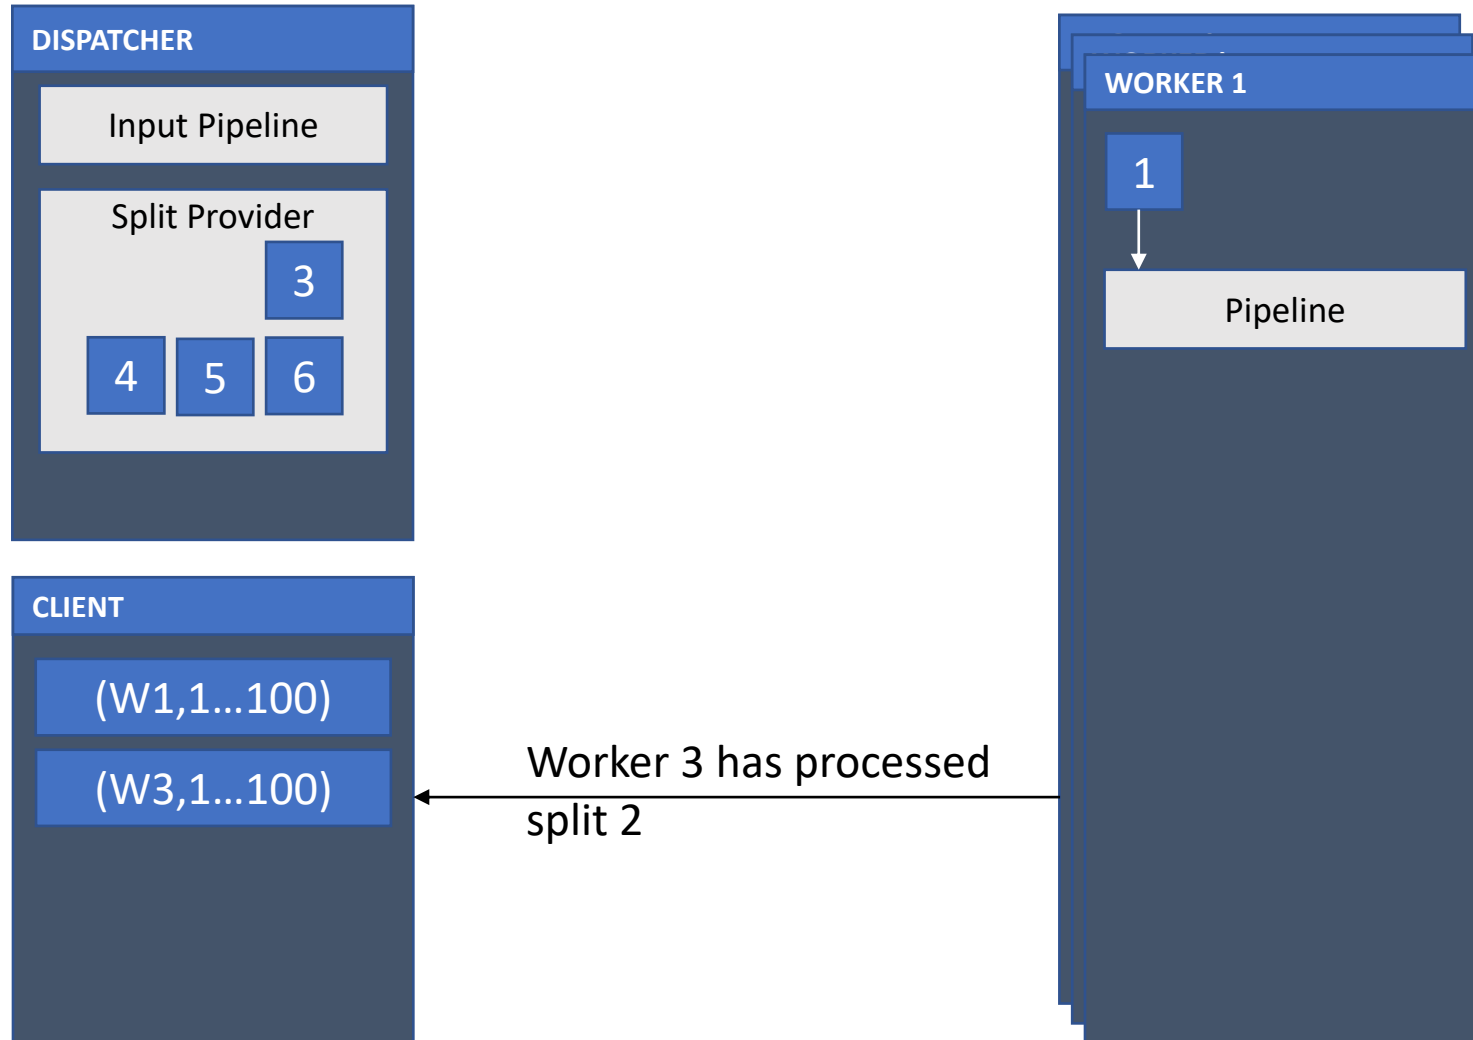
# Cachew Illustrated




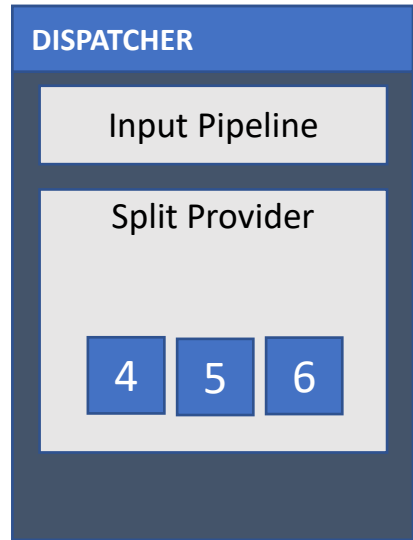
# Cachew Illustrated



# Cachew Illustrated


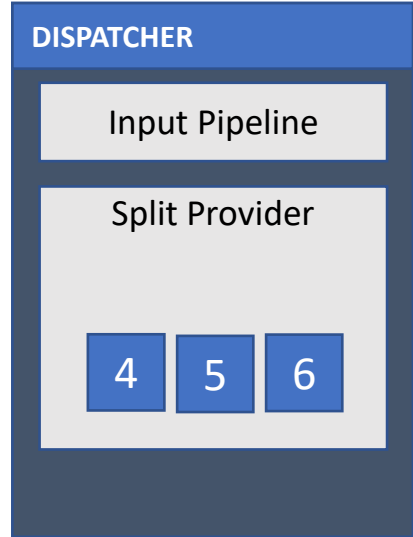


# Cachew Illustrated


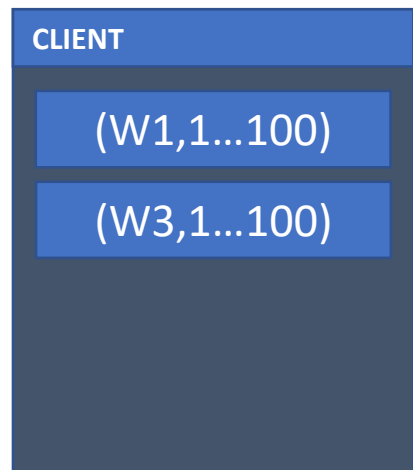
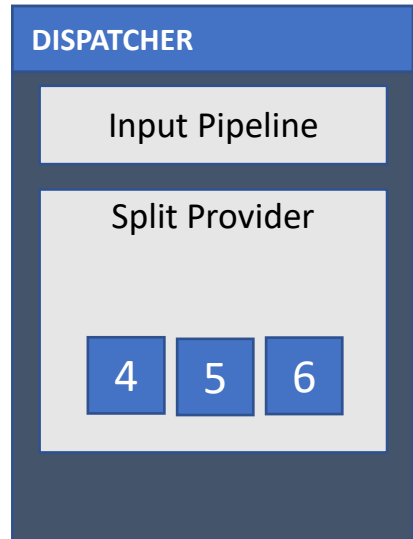




# Cachew Illustrated




# Cachew Illustrated



How can we react?

- Skip 3: at-most-once
- Recompute 1 & 3, transmit everything: at-least-once
- Recompute 1 & 3, but skip 1 at the client: exactly-once

# Cachew Illustrated




How can we react?

- Skip 3: at-most-once
- Recompute 1 & 3, transmit everything: at-least-once
- Recompute 1 & 3, but skip 1 at the client: exactly-once


What do we need?

- Dispatcher: Remember split <-> worker association
- Client: Remember expected Batch ID
- Detect failure, failover task to a different worker
- Deterministic ordering

# Failover Illustrated



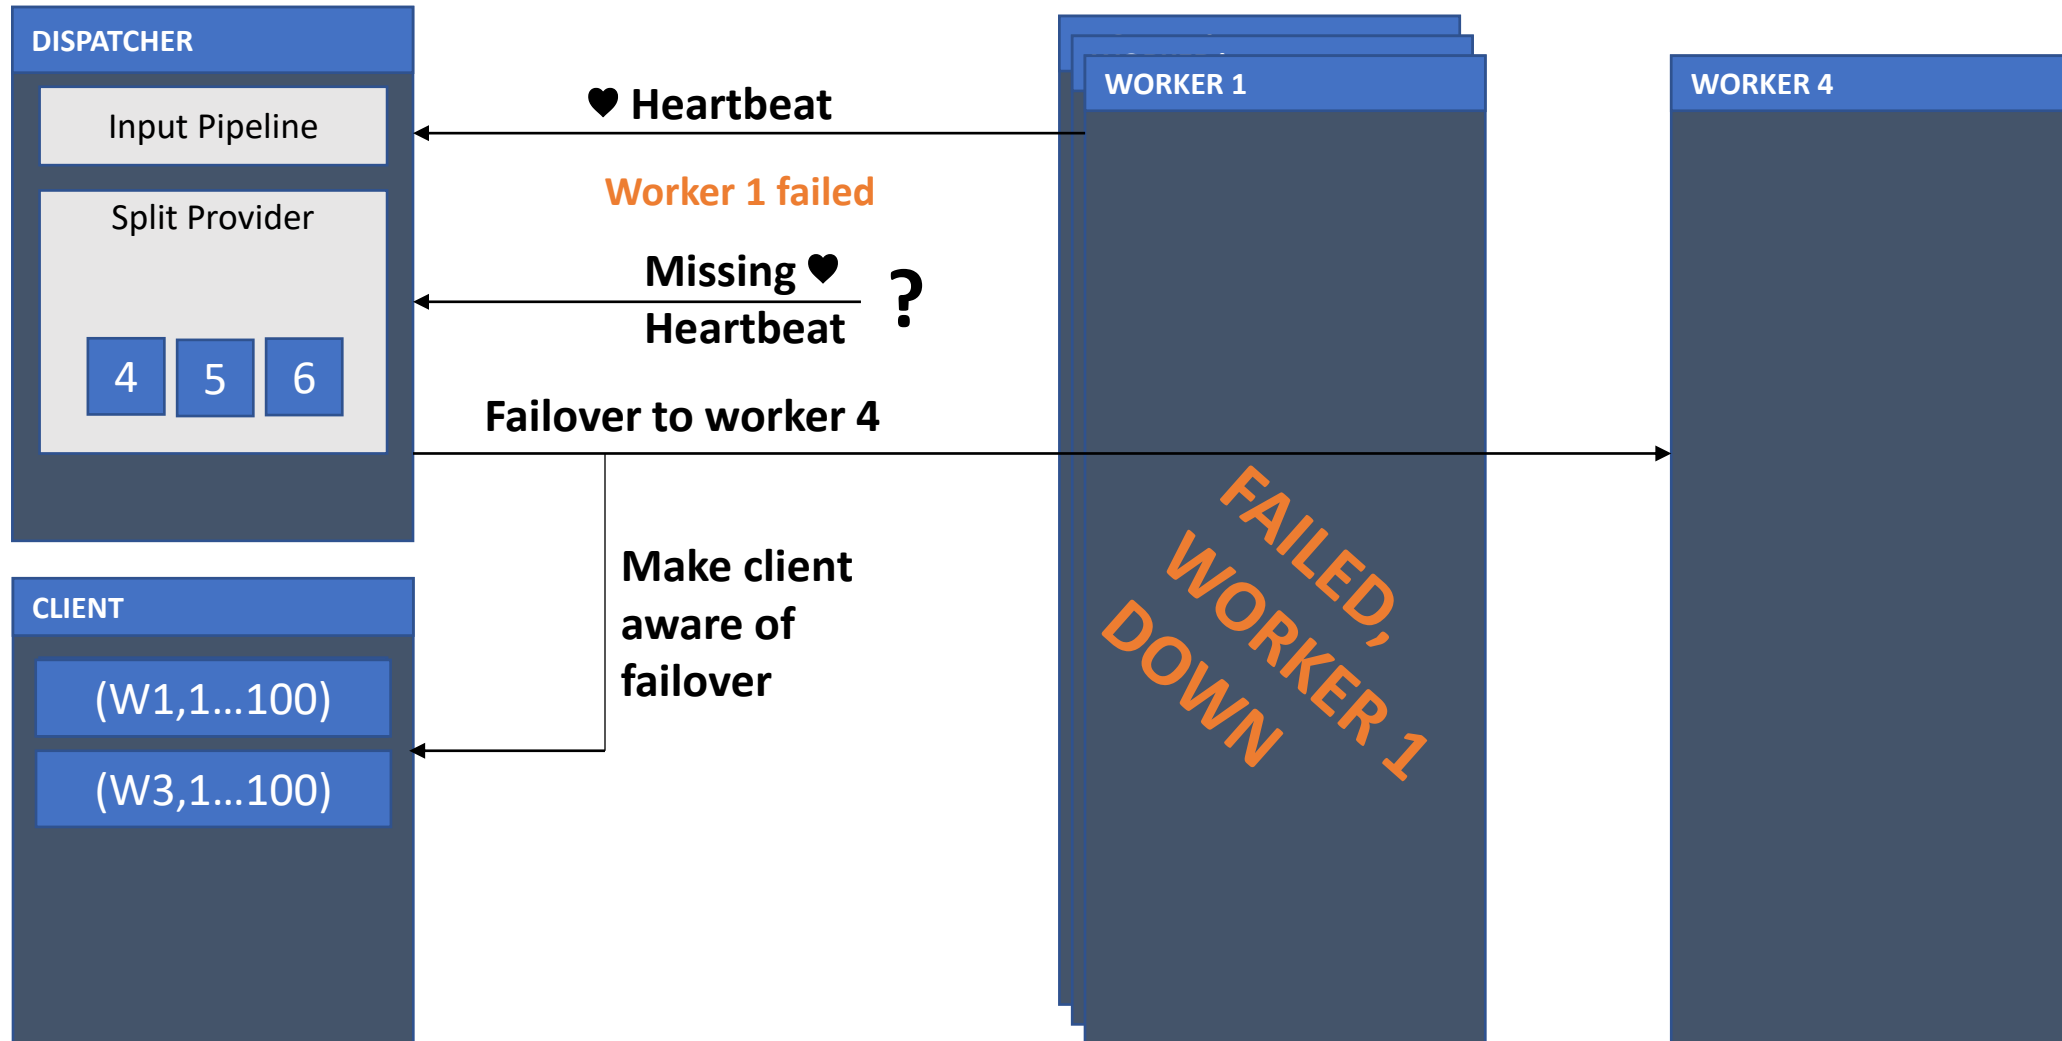
# Failover Illustrated



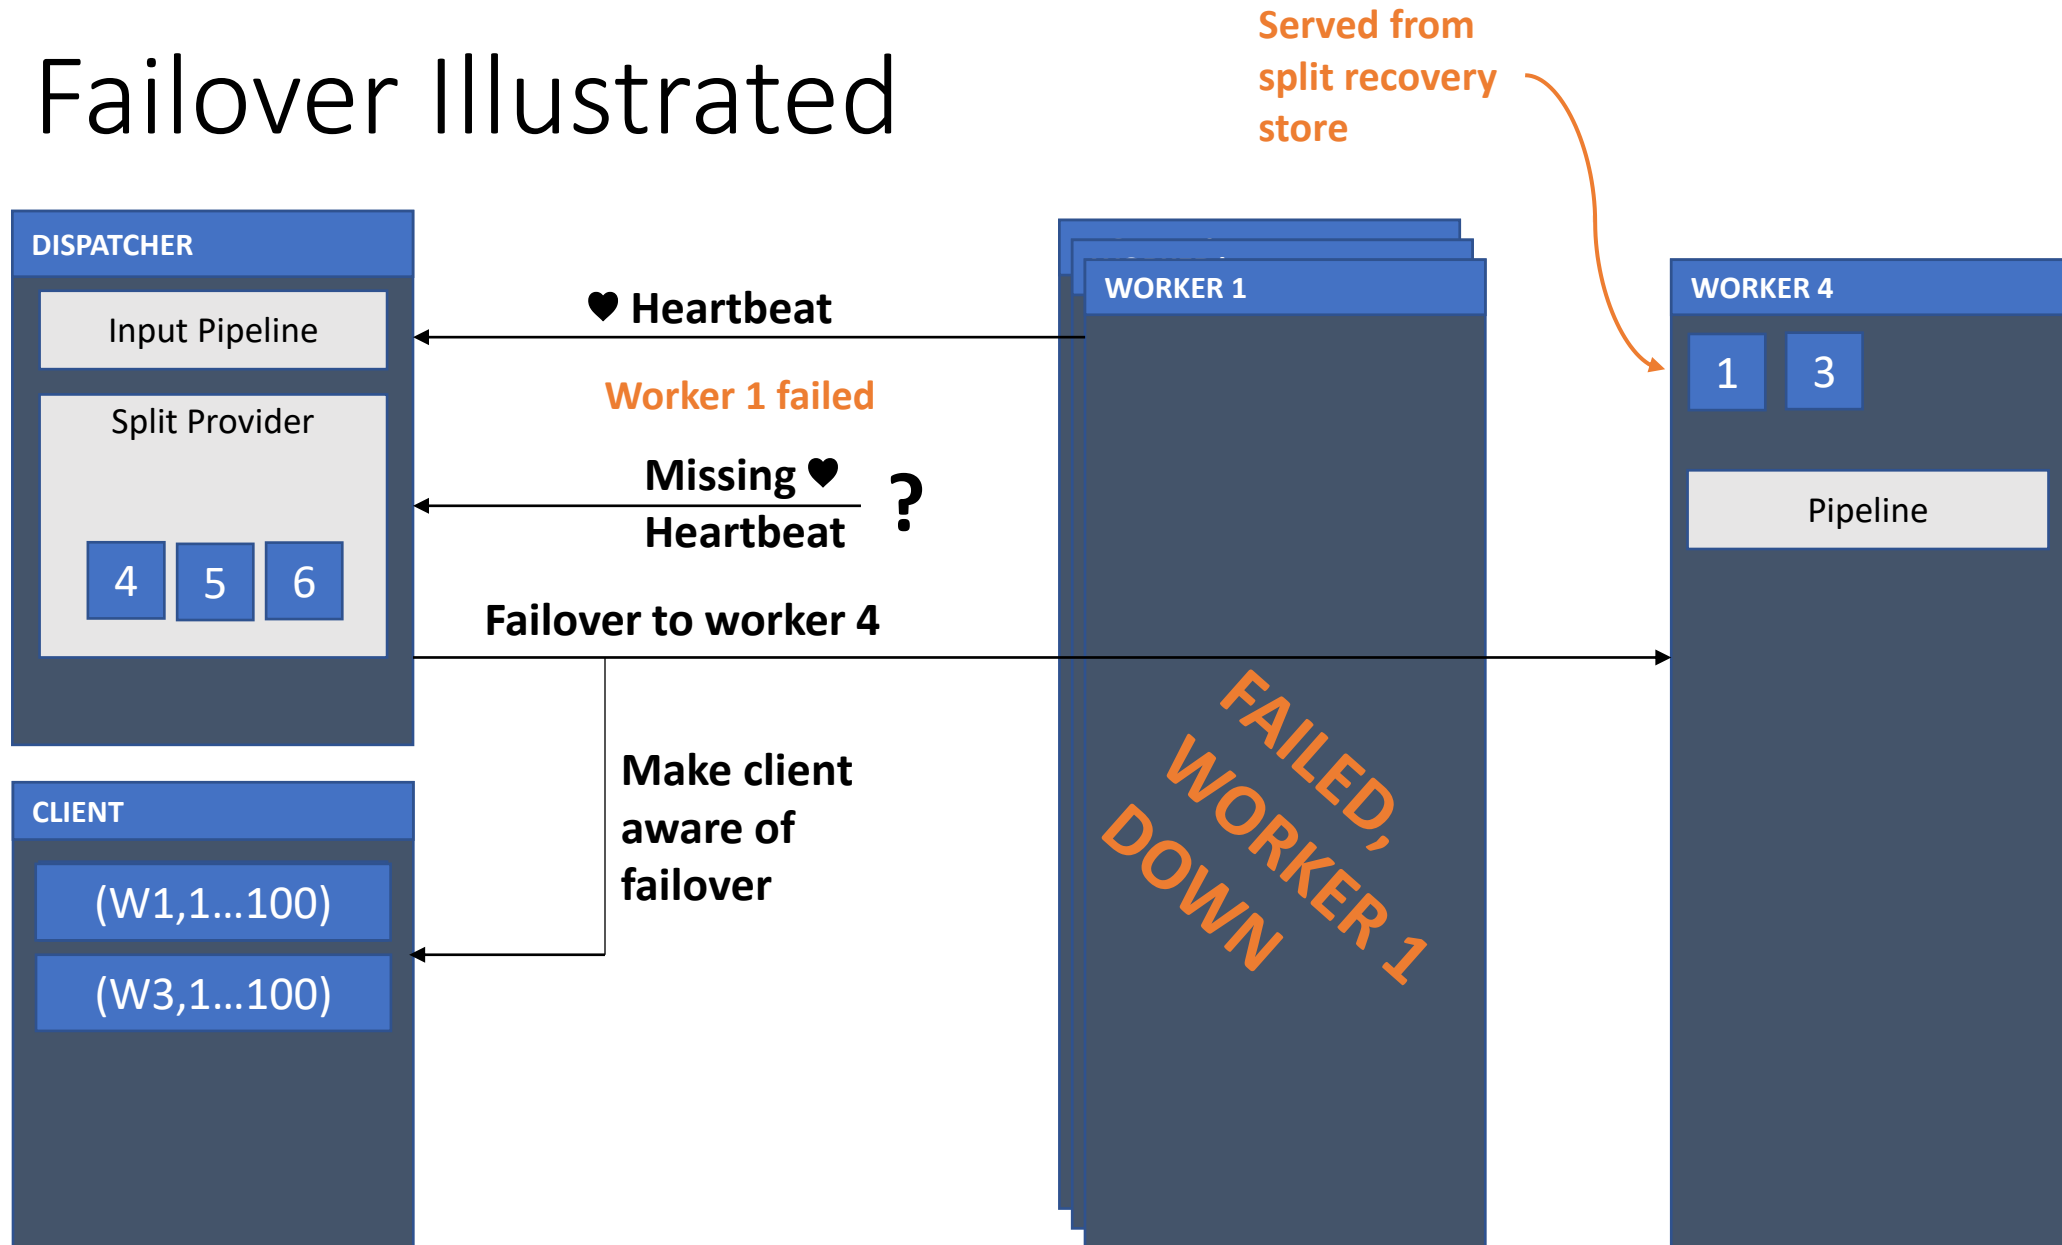
# Failover Illustrated



# Failover Illustrated

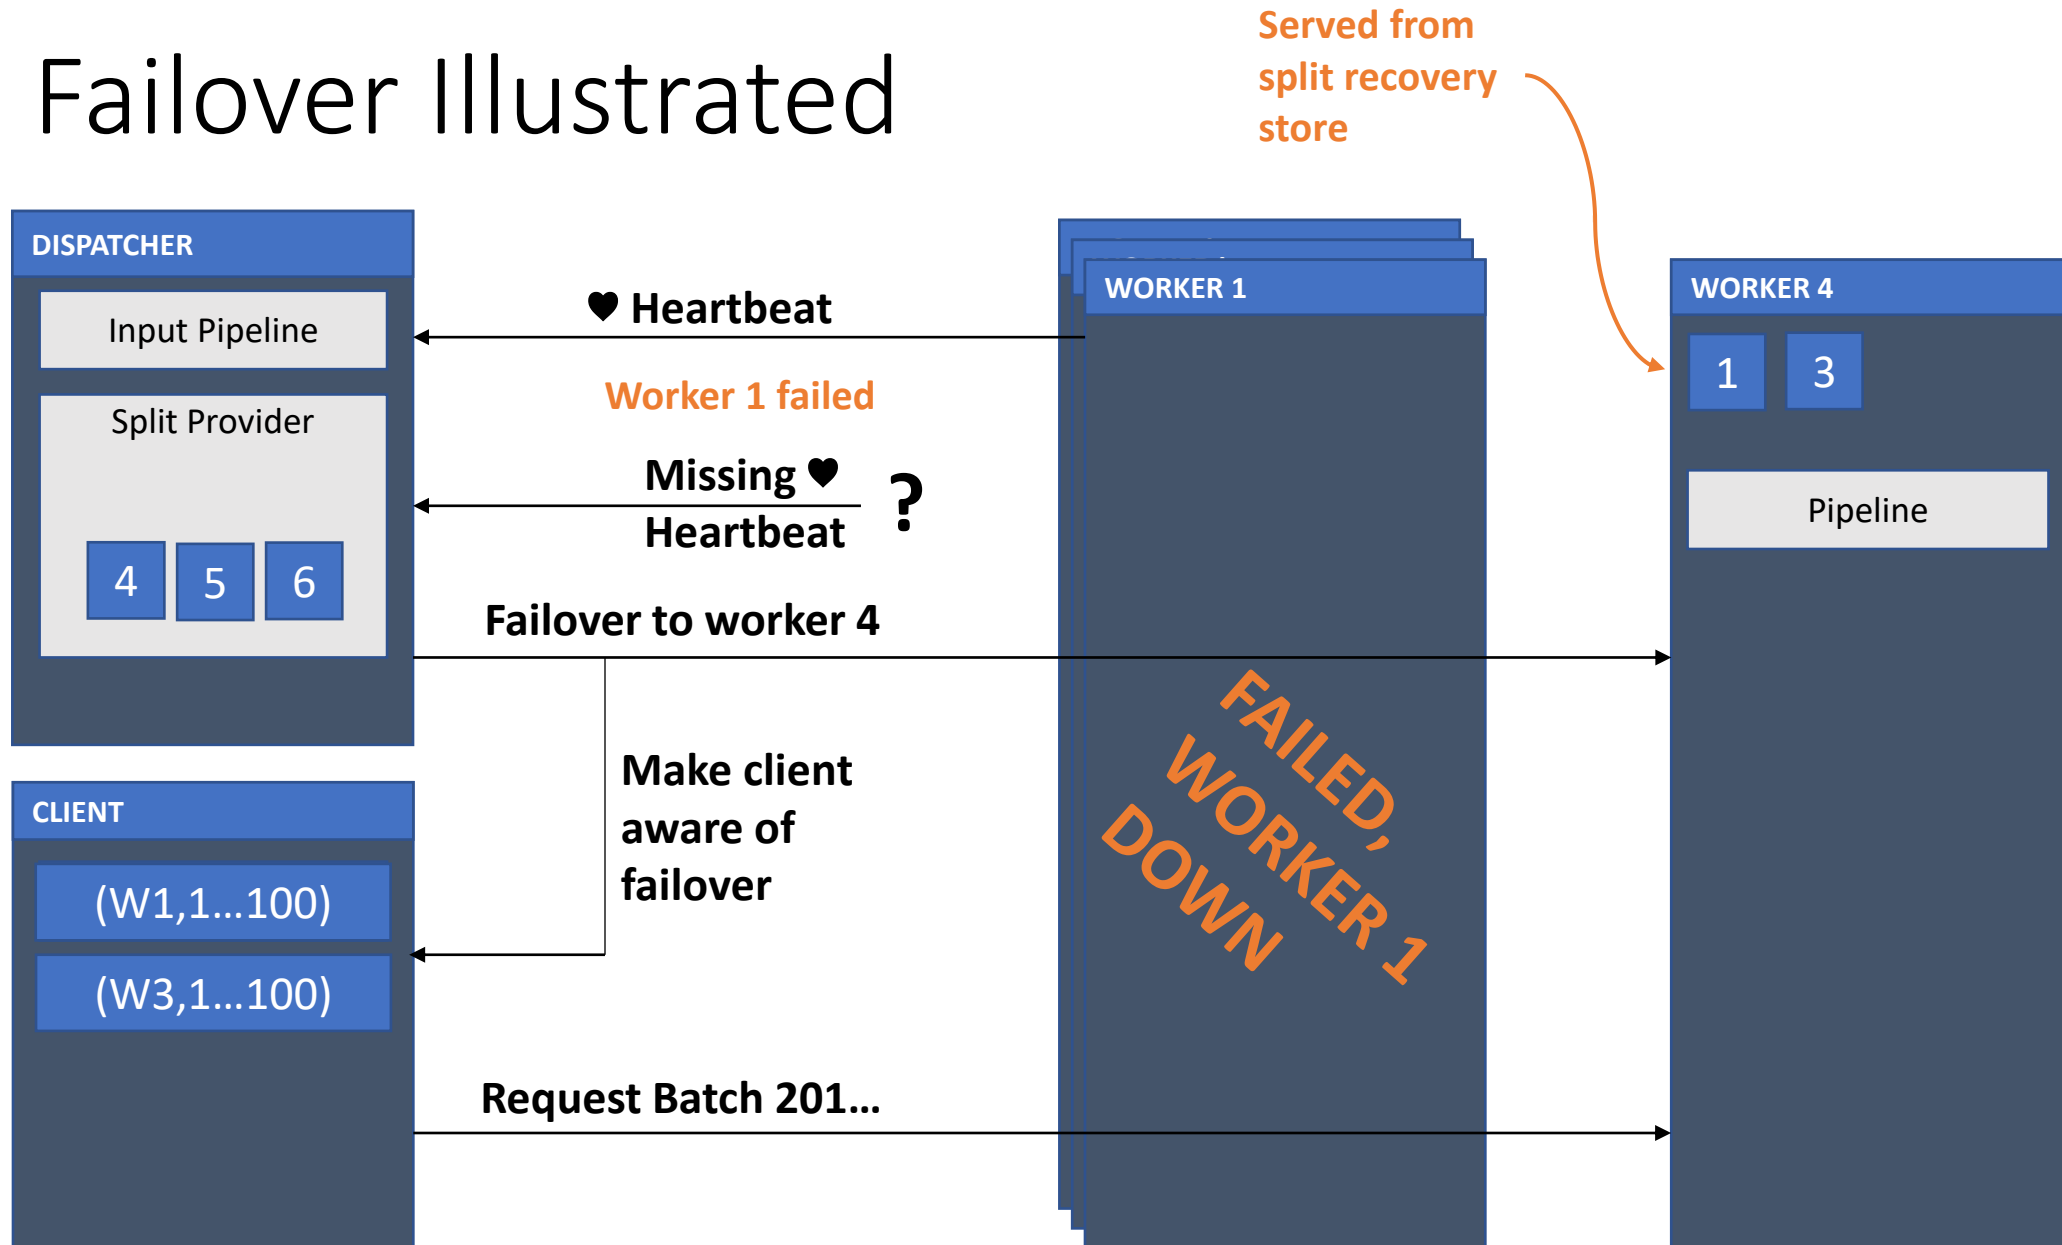


# Failover Illustrated

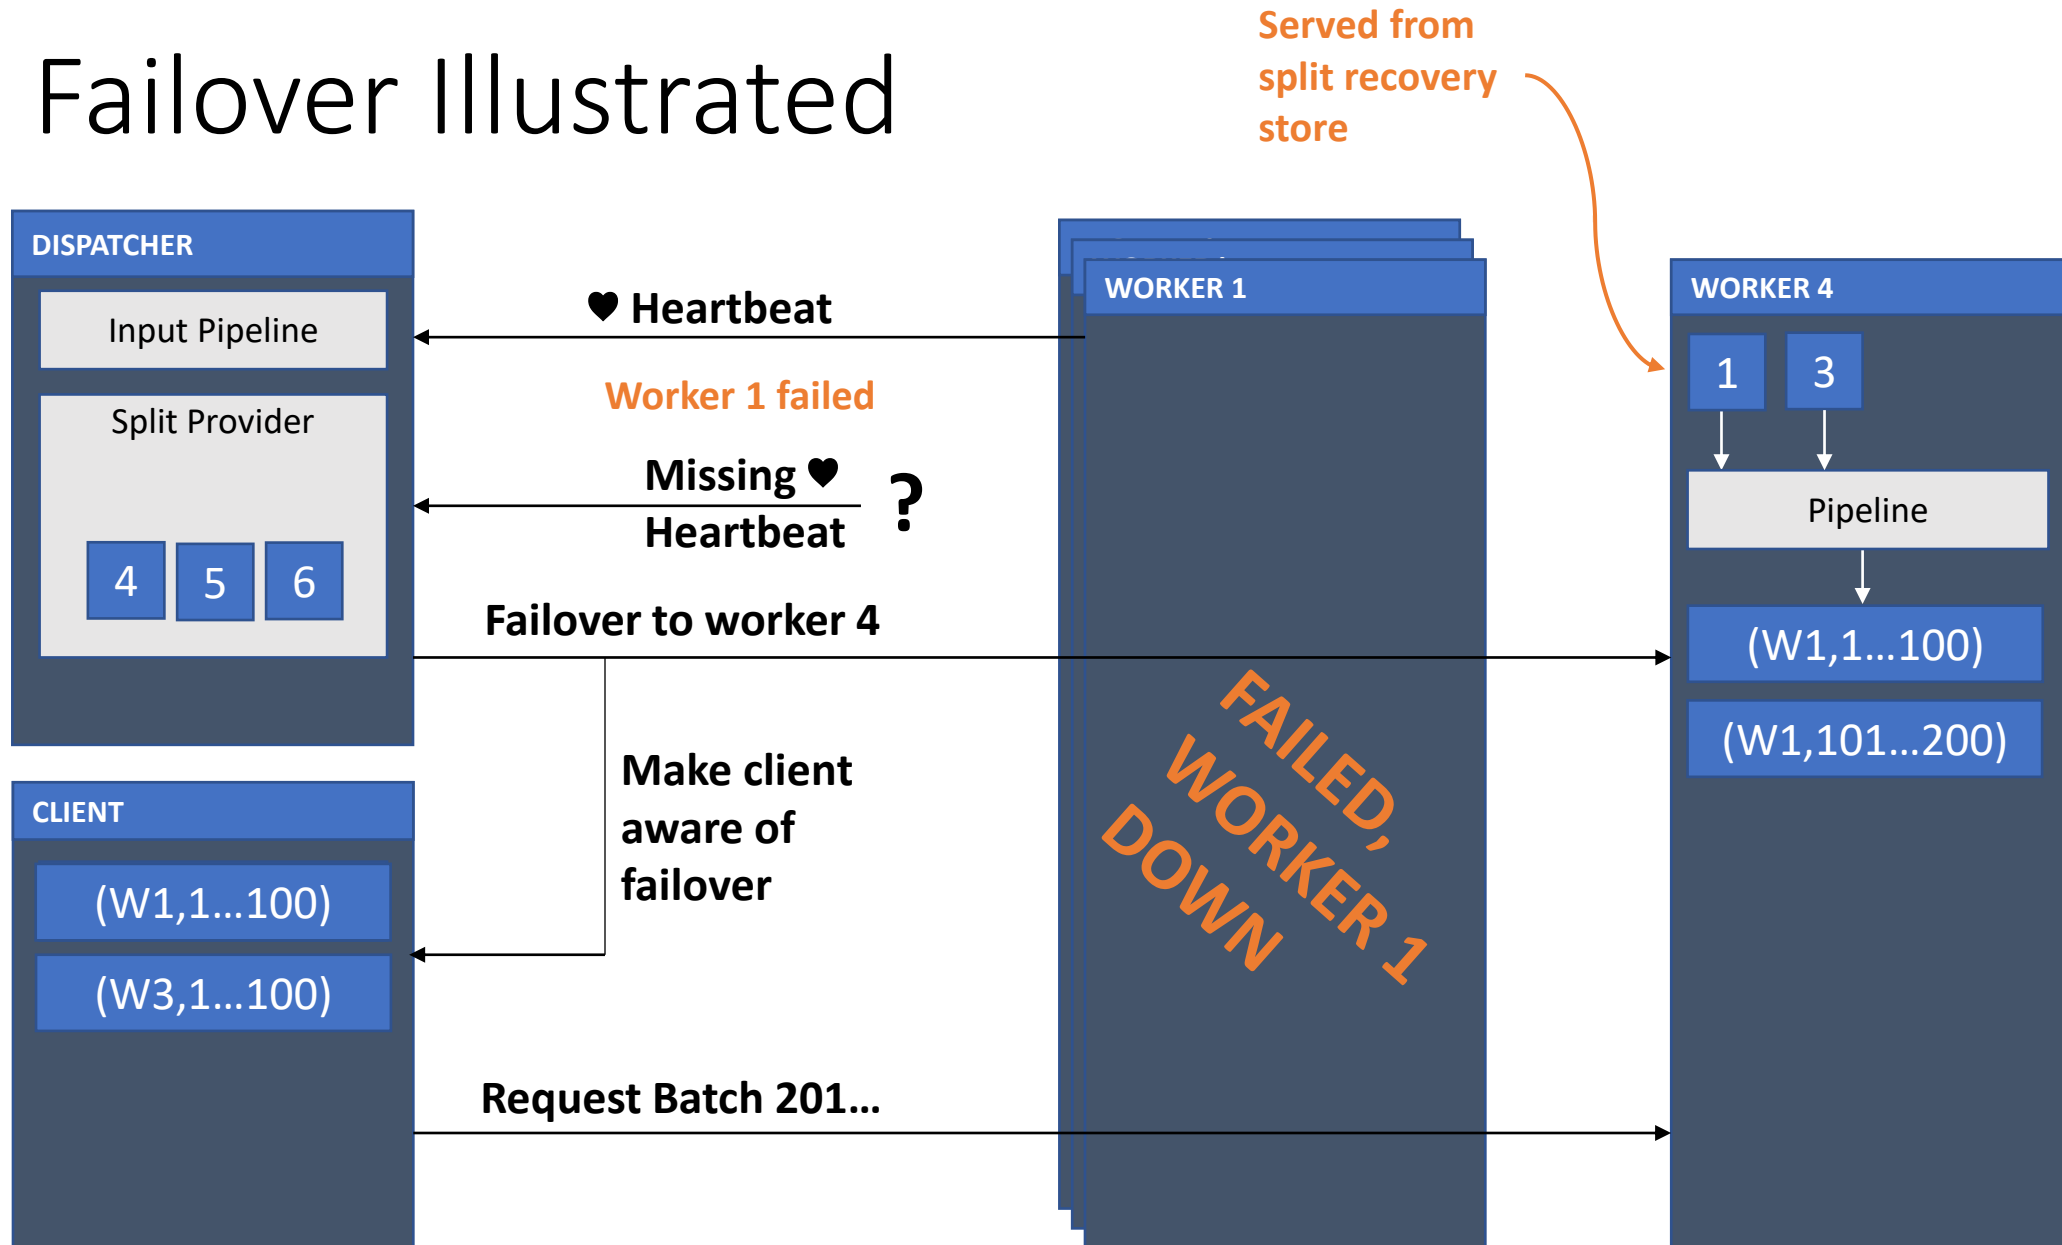




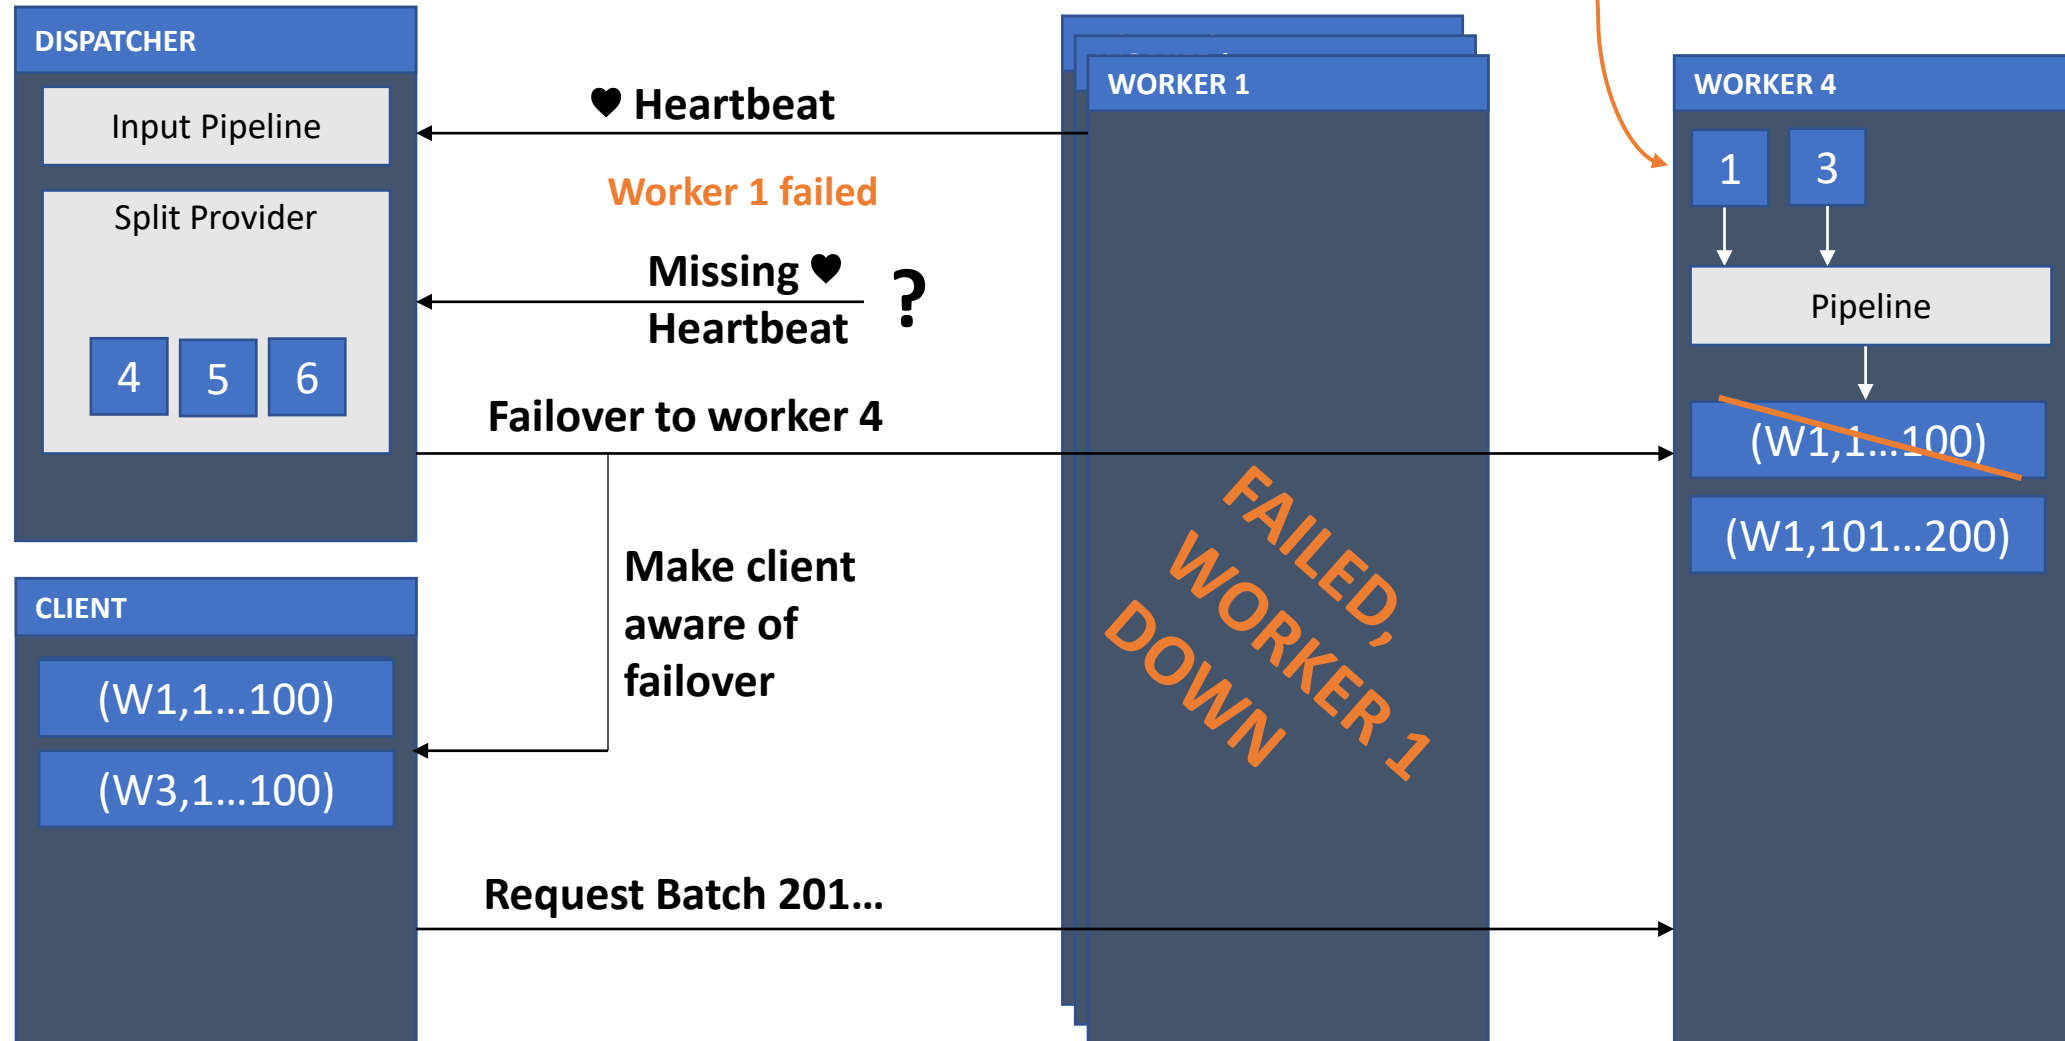
# Failover Illustrated



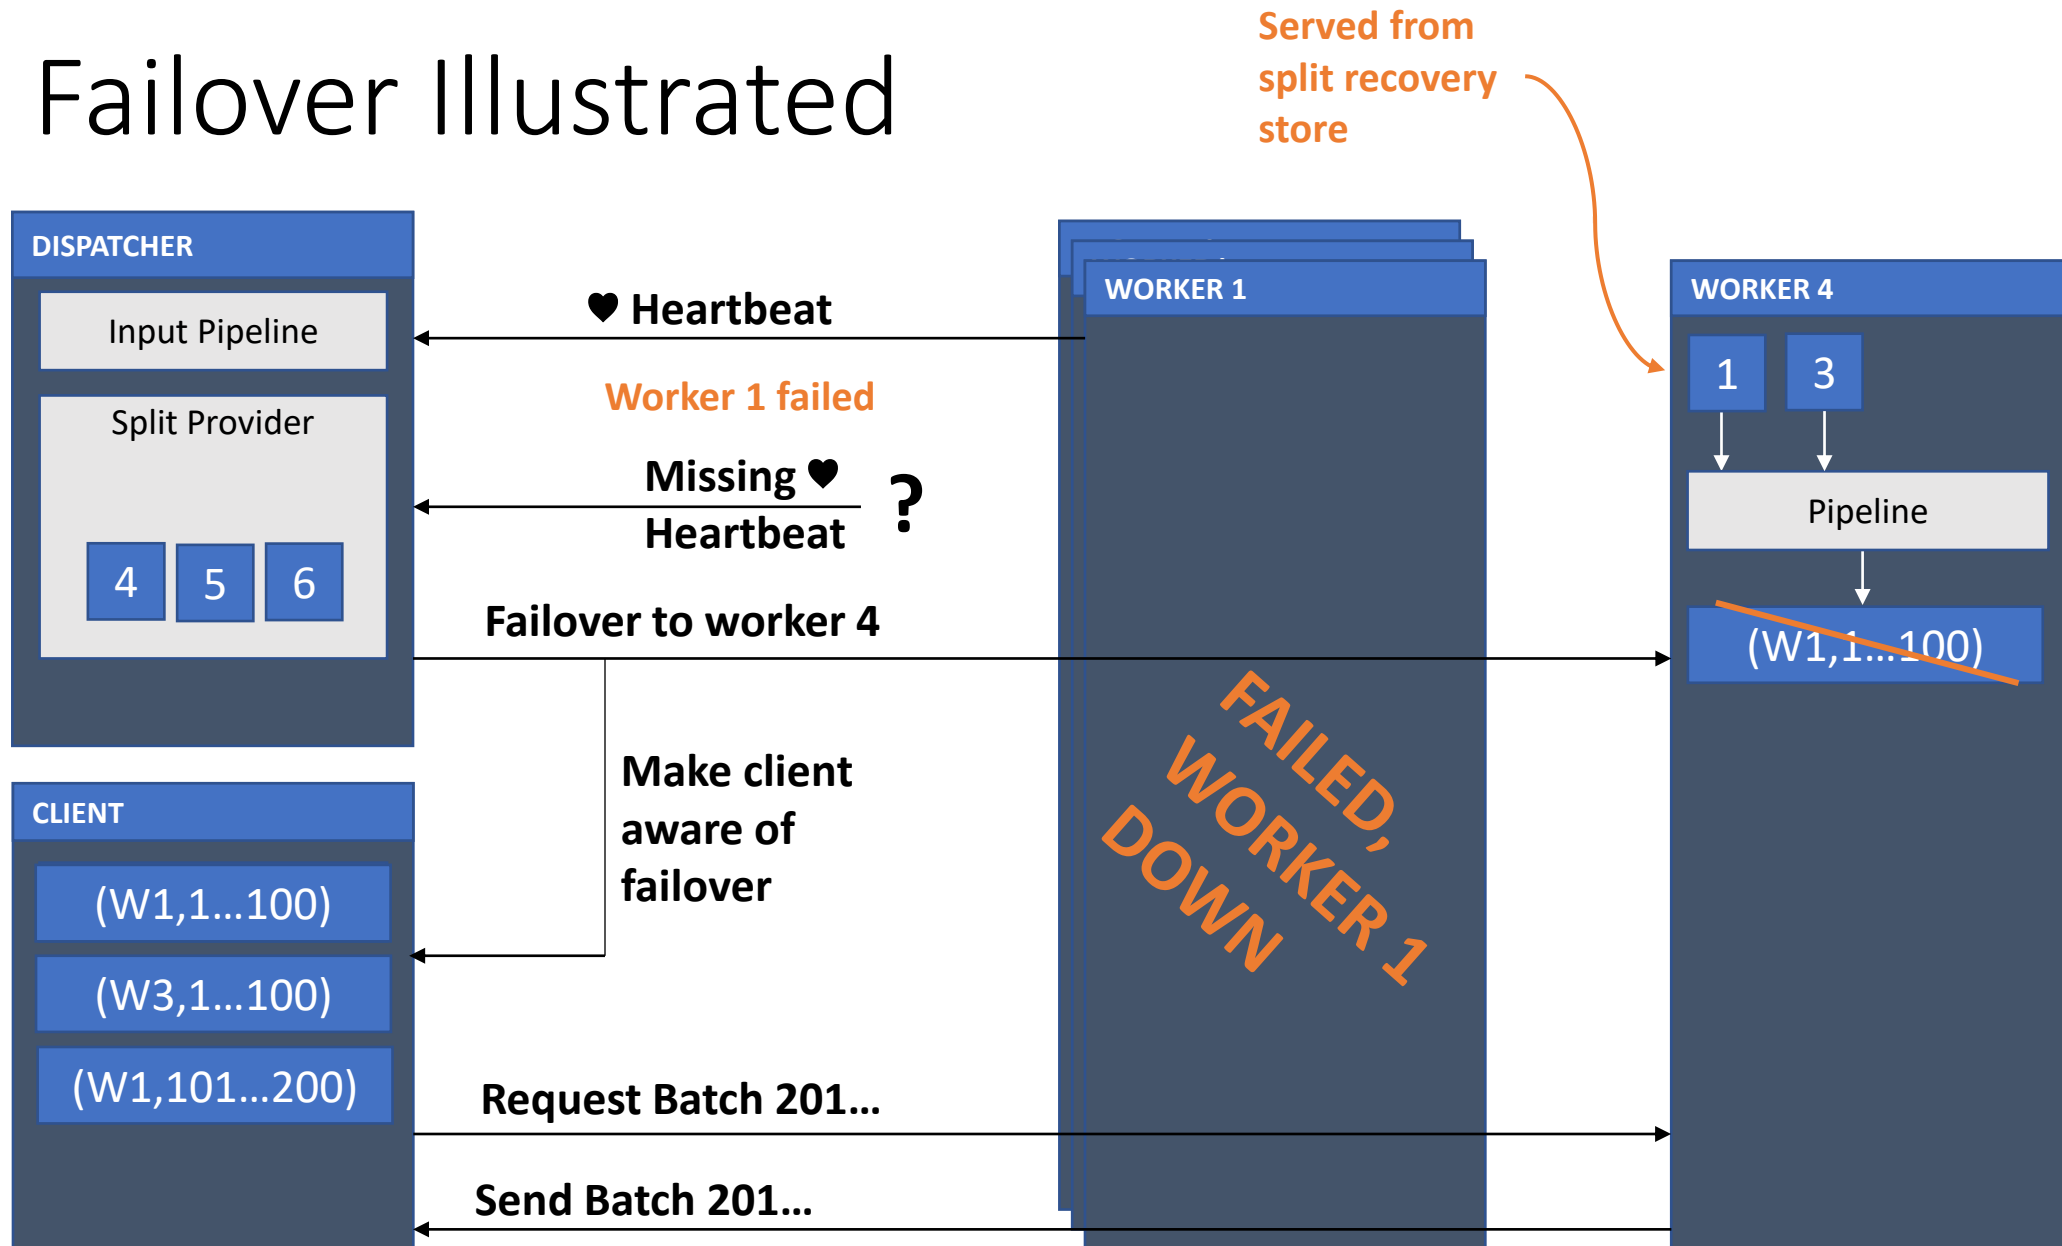
# Failover Illustrated



# Failover Illustrated



# Failover Illustrated



# Reducing Overhead

overhead = detection + failover + recomputation

## Detection

Fine-tune heartbeat  
detection mechanisms

~ 1-10s

## Failover

Have nodes in hot-standby  
Implement “express”  
messages (circumvent  
heartbeat-based protocols)

~ 1-10s

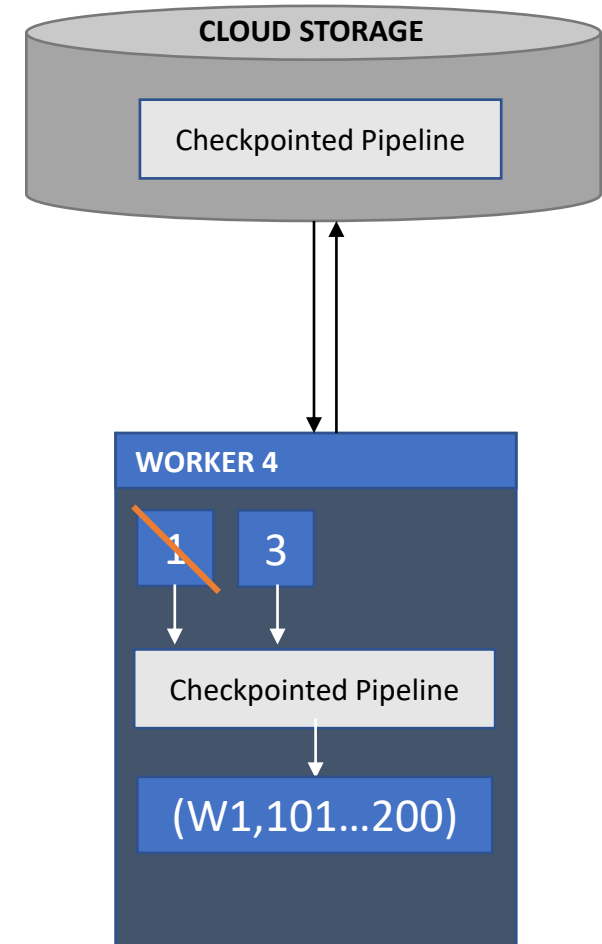
## Recomputation


Checkpoint worker state to  
recompute less

~ ½ epoch

# Checkpointing

- Workers regularly checkpoint their pipeline state to GlusterFS (cloud storage)
- This reduces the overhead because we need not recompute everything from scratch (Req3)
- The pipeline needs to be locked to arrive at a consistent state (checkpoint stall)
- Implementation builds on top of existing tf.data checkpointing mechanism






## Takeaways

- 1) Throughput is constant; important characteristic used for scaling/profiling
- 2) Checkpointing introduced overhead

Throughput under Worker Failure  
[Batches / Second]



# Takeaways

- 1) Without any worker checkpointing we have more recomputation
- 2) Worker checkpoints recover some of the progress (Req3)




# Overhead Revisited

overhead = detection + failover + **checkpointing** + recomputation

## Checkpointing Overhead

- Pipeline needs to be checkpointed in a consistent state (so it is locked)
- This will result in a drop in throughput
- Overhead is highly dependent on the structure and contents of individual pipelines
- Tradeoff between checkpointing and recomputation overhead
- Depends highly on the failure pattern

Overhead of a Single Checkpoint  
epoch time increase [multiplicative]



# Questions

How does checkpointing overhead compare for different pipelines?

Is it feasible for all pipelines to make even just a single checkpoint?

# Takeaways

1) Overhead varies with pipeline

2) Sometimes no checkpoint is better

# Conclusion

- Built a system which satisfies all requirements (distributed, exactly-once, performant, bounded overhead and reproducible randomness) in the C++ layer maintaining full compatibility with all tf.data input pipelines
- Evaluation shows that worker checkpoints are not always feasible

# Further Work

- Optimal checkpointing frequency
- Optimize checkpointing overhead (e.g. specific Ops)
- Use small independent sets of work (like Meta's DPP)
- Try to flush out pipelines (as in Meta's Check-N-Run)

# References

- Jayashree Mohan, Amar Phanishayee, and Vijay Chidambaram. Check-Freq: Frequent, Fine-Grained DNN checkpointing. In 19th USENIX Conference on File and Storage Technologies (FAST 21), pages 203–216. USENIX Association, February 2021.
- Derek G. Murray, Jiří Šimša, Ana Klimovic, and Ihor Indyk. Tf.data: A machine learning data processing framework. Proc. VLDB Endow.,
- JCS Kadupitige, Vikram Jadhao, and Prateek Sharma. Modeling the temporally constrained preemptions of transient cloud vms. In Proceedings of the 29th International Symposium on High-Performance Parallel and Distributed Computing, HPDC '20, page 41–52, New York, NY, USA, 2020. Association for Computing Machinery.
- Dan Graur, Damien Aymon, Dan Kluser, Tanguy Albrici, Chandramohan A Thekkath, and Ana Klimovic. Cachew: Machine learning input data processing as a service. In 2022 USENIX Annual Technical Conference (USENIX ATC 22), pages 689–706, 2022.
- Assaf Eisenman, Kiran Kumar Matam, Steven Ingram, Dheevatsa Mudigere, Raghuraman Krishnamoorthi, Krishnakumar Nair, Misha Smelyanskiy, and Murali Annavaram. {Check-N-Run}: a checkpointing system for training deep learning recommendation models. In 19th USENIX Symposium on Networked Systems Design and Implementation (NSDI 22), pages 929–943, 2022.

## Convolutud Pipeline Example

